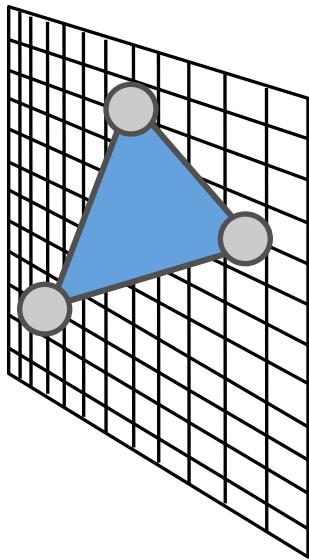


# Shadows

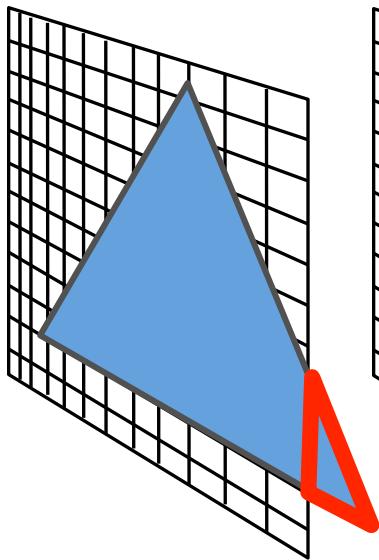
©Anthony Steed 1999-2003, 2014, Yiorgos Chrysanthou 2013, Tobias Ritschel 2016



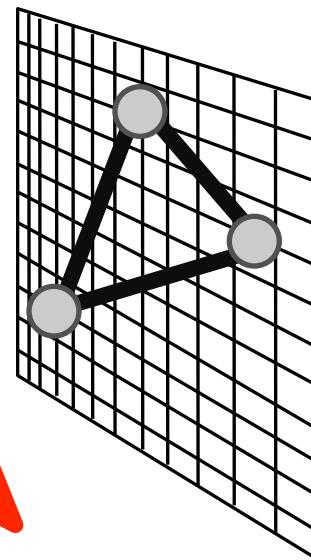
# Challenges



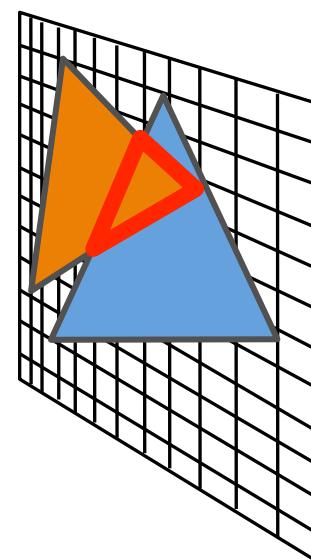
Projection



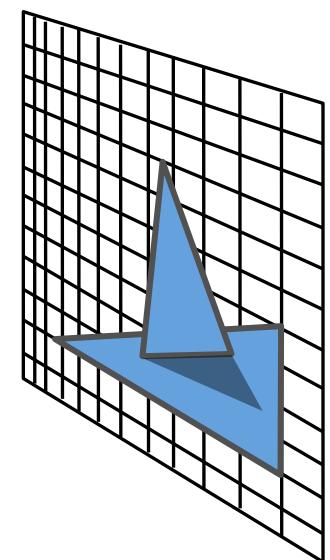
Clipping



Rasterization

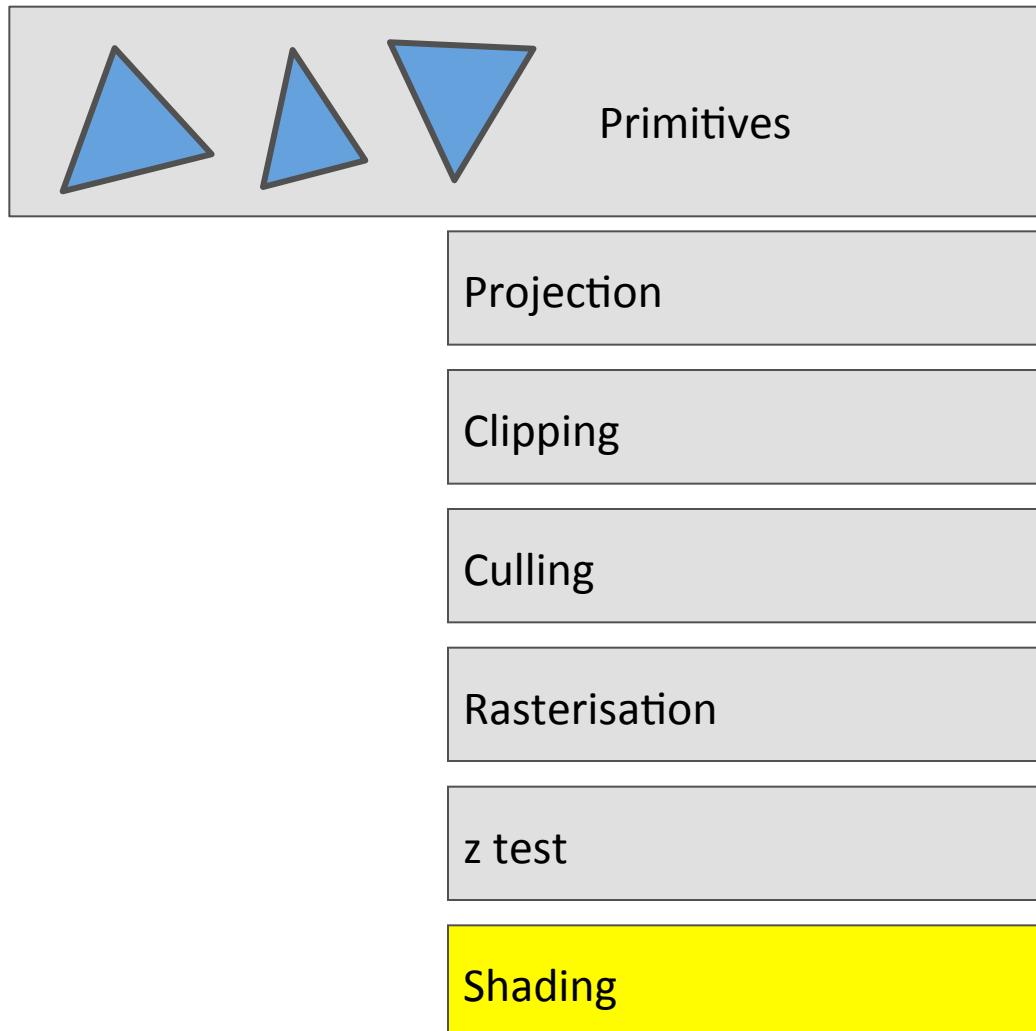


Visibility



Shading

# Pipeline

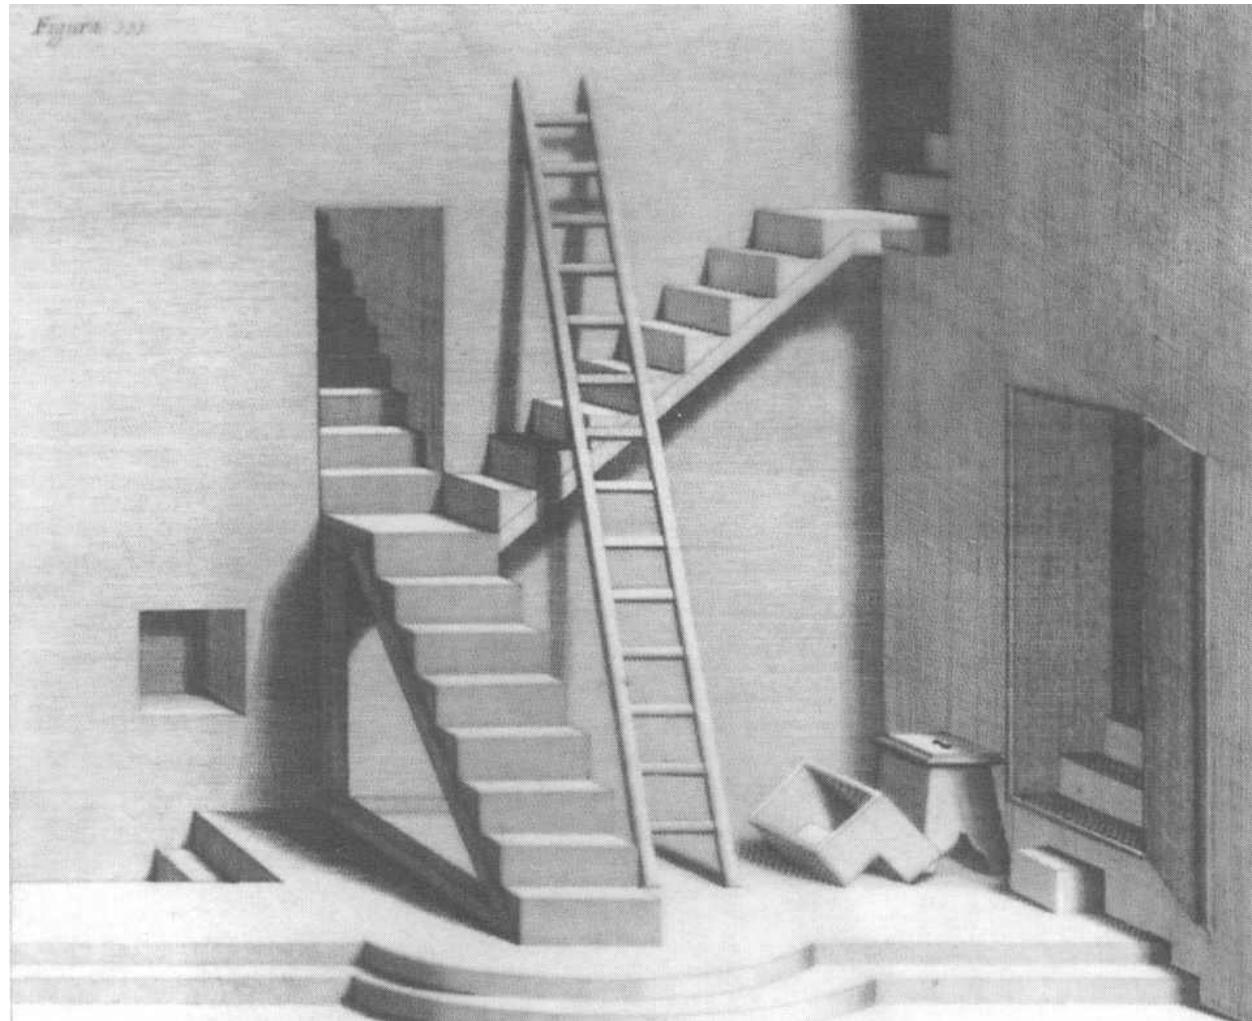


# Outline

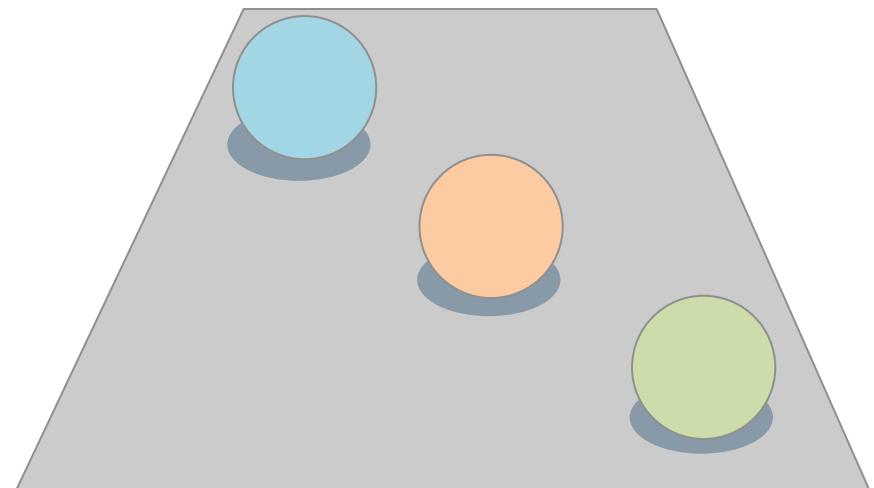
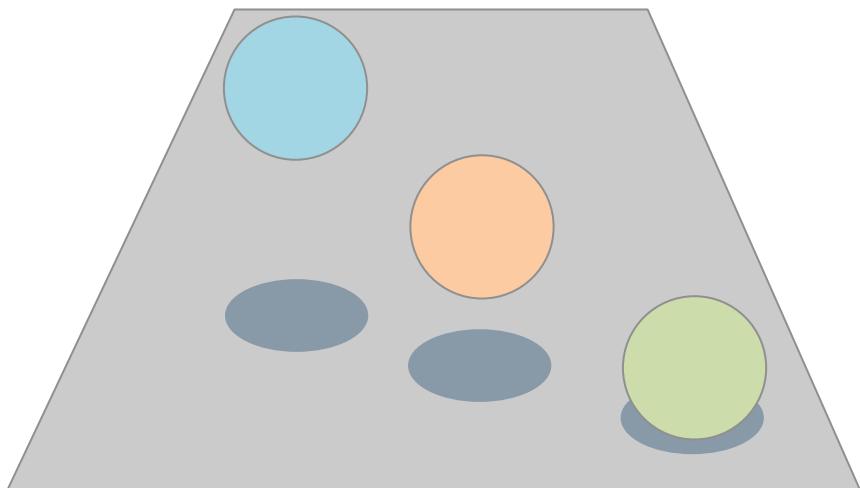
- Introduction
- Sharp shadows
- Soft shadows
- Conclusion

# Why are Shadows Important?

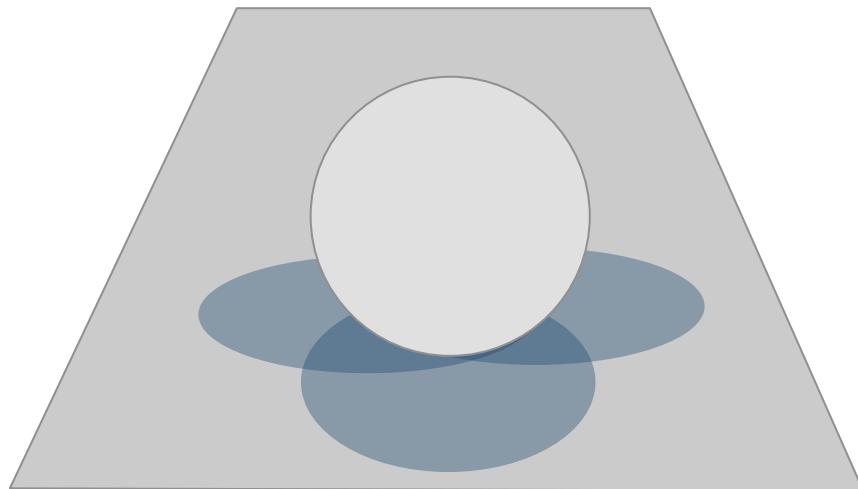
- Depth cue
- Scene
- Lighting
- Realism
- Contact points



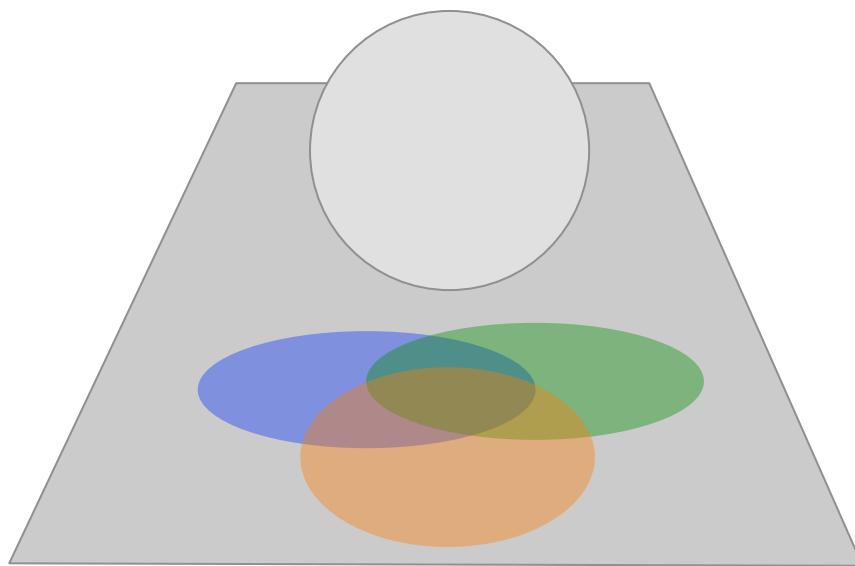
# Spatial cue



# Light position cue (sundial)

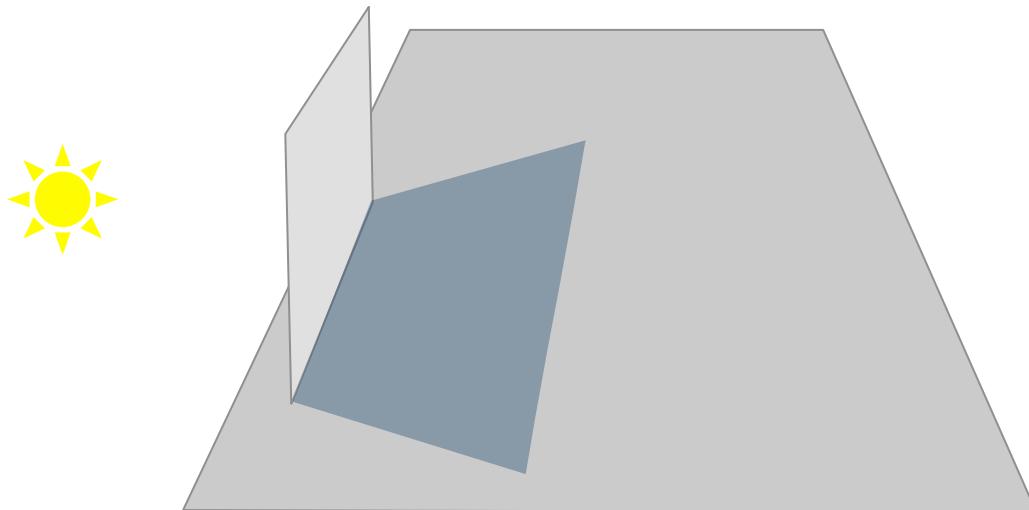


# Light color

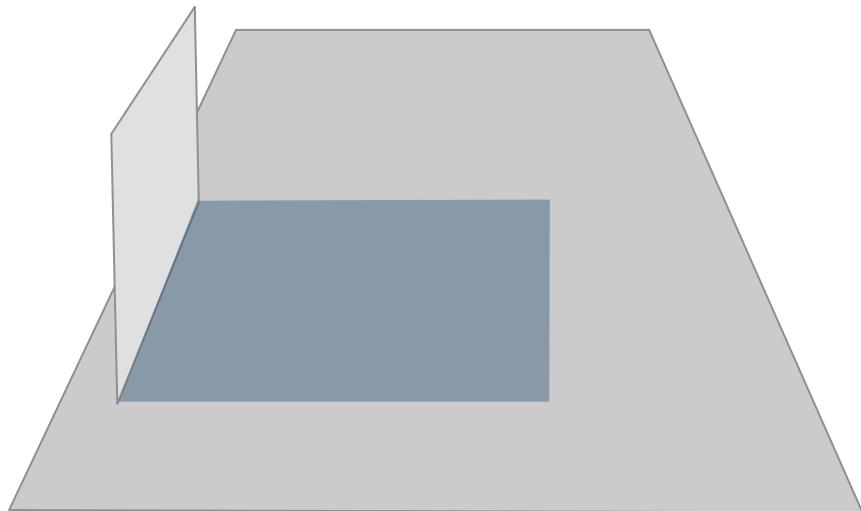




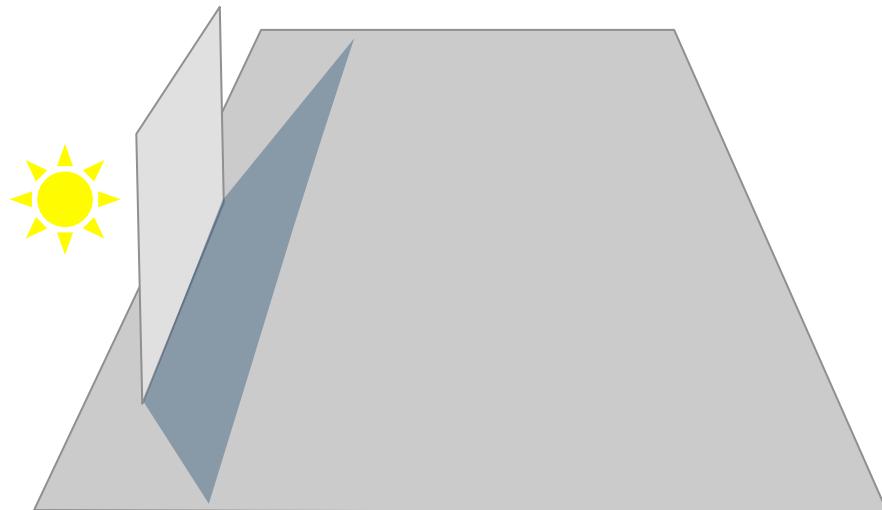
# Light distance



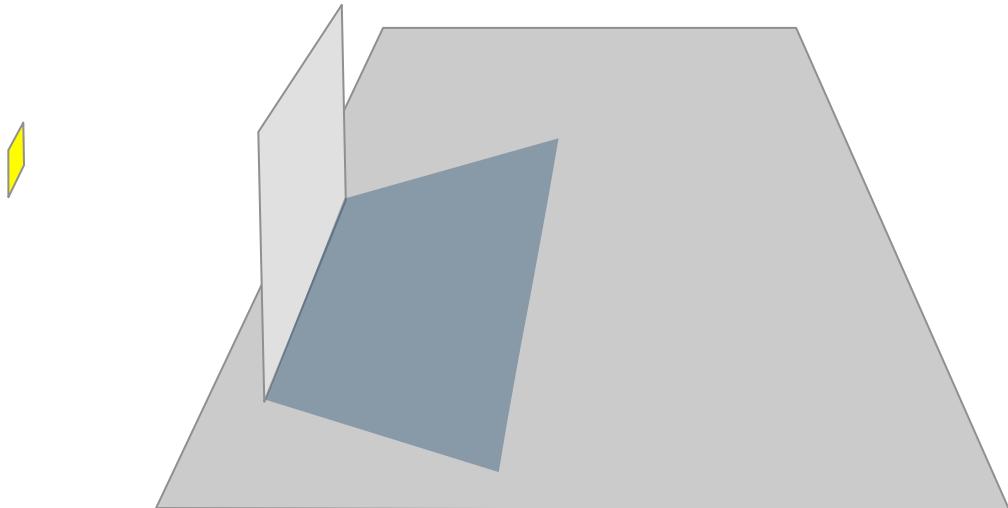
# Light distance (far)



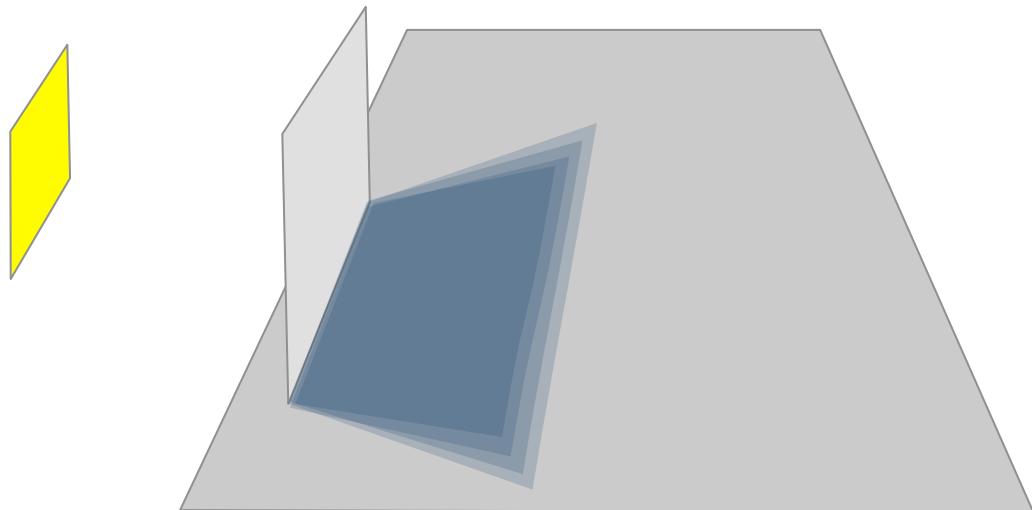
# Light distance (near)



# Light size (small)



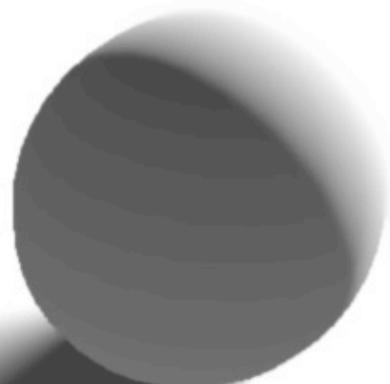
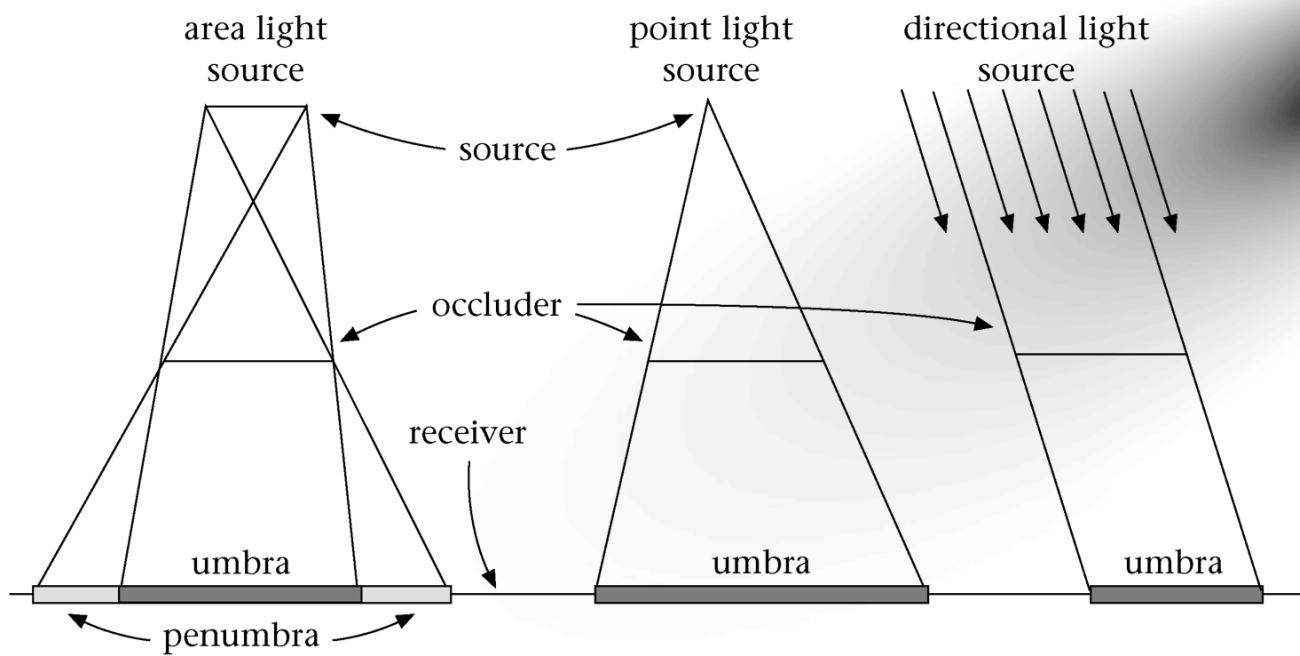
# Light size: large



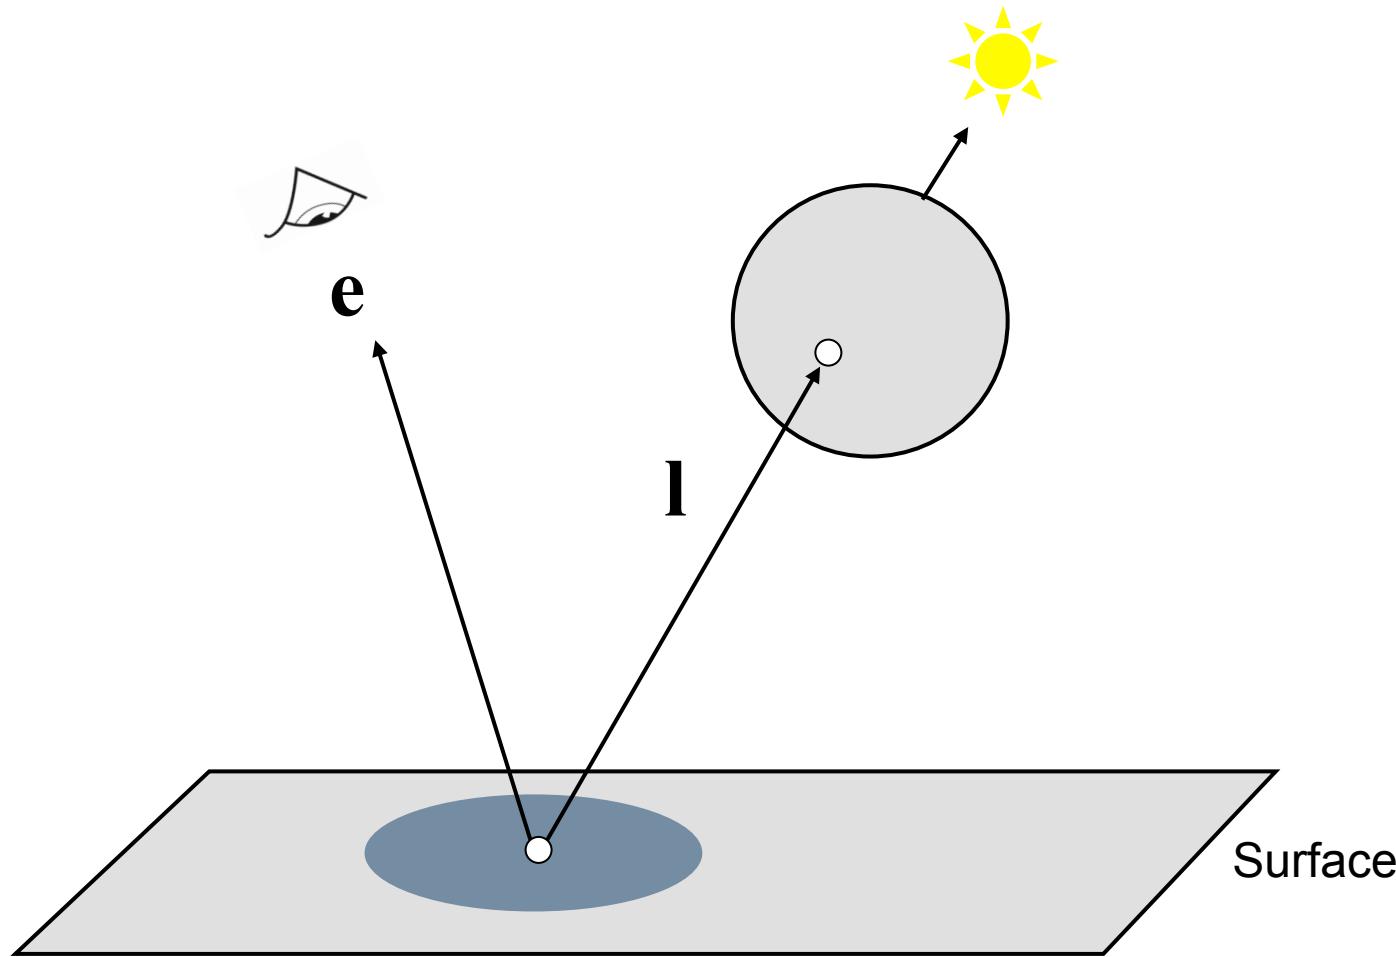
# Shadows are Complex

- In the real world sources of light are not points
- The intensity within a shadow is not constant
  - **Umbra**  
part that sees nothing of the source
  - **Penumbra**  
part that receives some light

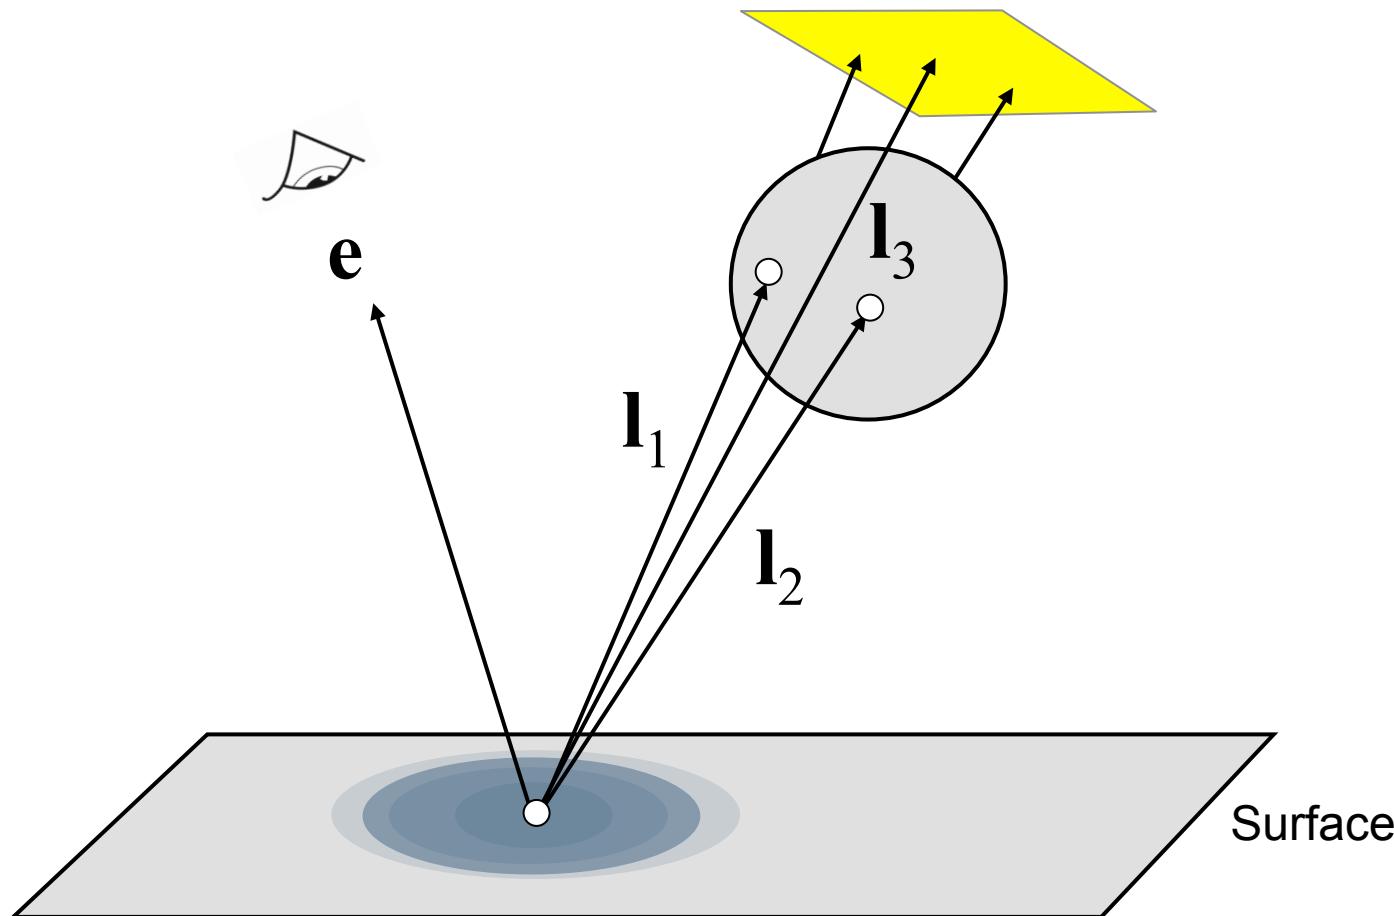
# Definition



# Shadow in ray-tracing

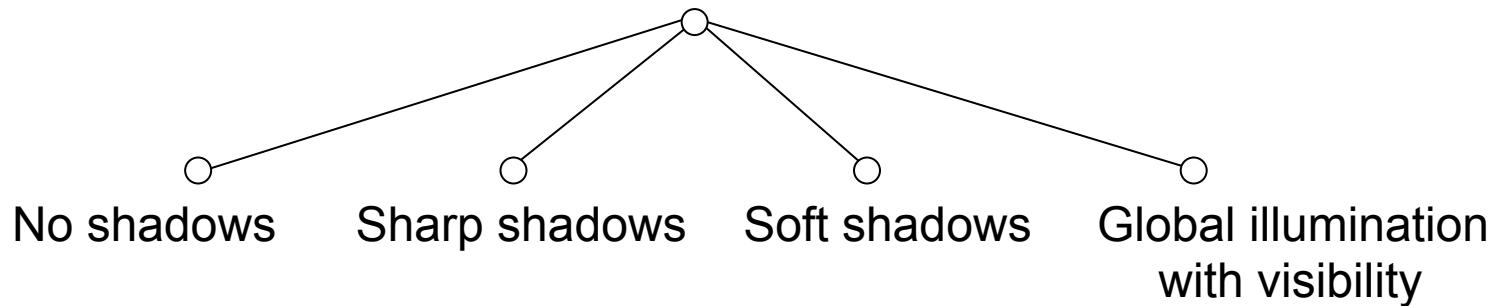


# Soft shadows in ray-tracing



# Current Shadowing Methods

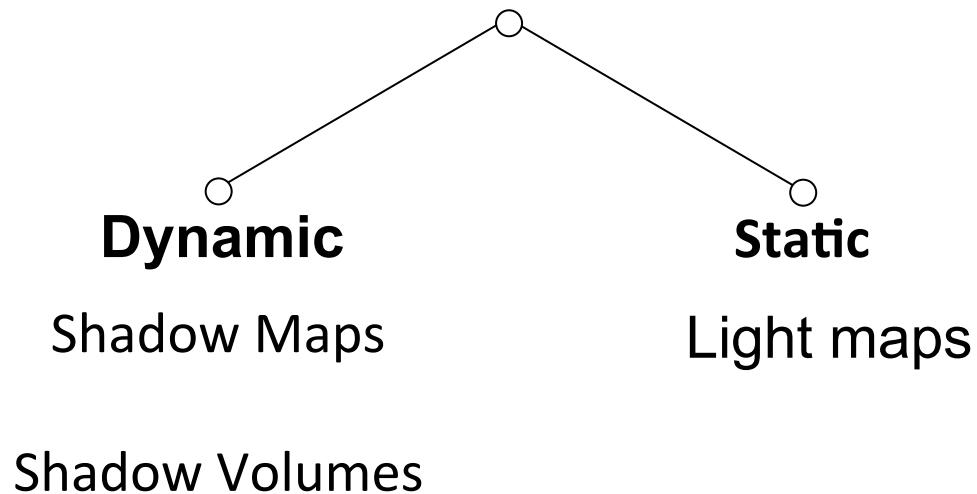
- There exist a large number of methods
- We are interested in methods suitable for interactive walkthroughs, speed is crucial
- We will classify them by complexity:



# Sharp shadows

# Sharp Shadows

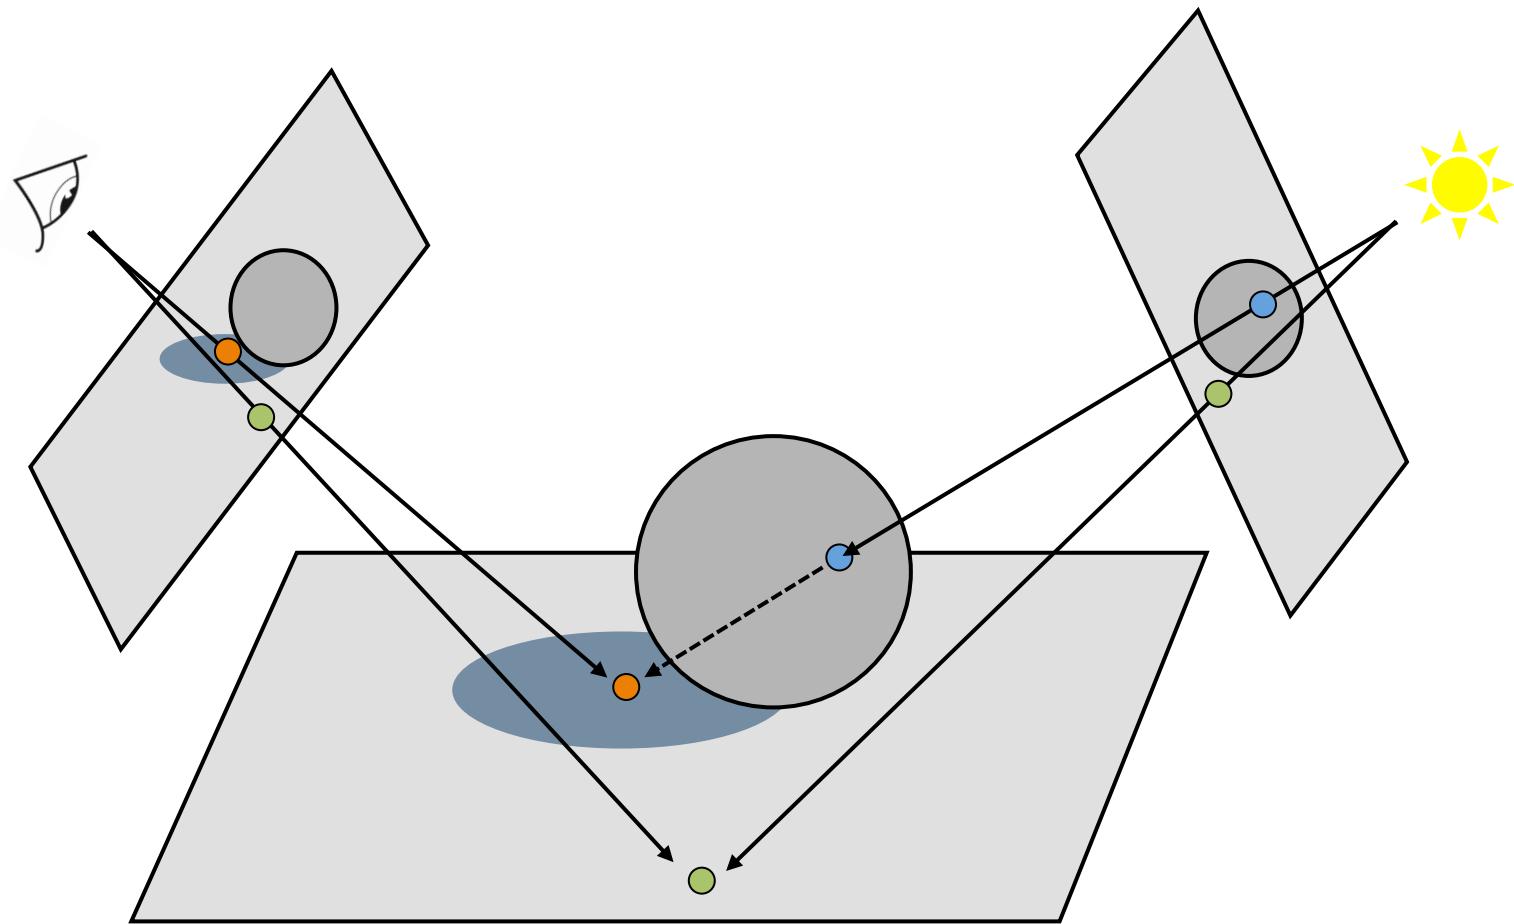
Source is **assumed** to be a point or direction



# Shadow Maps

Light / shadow duality:

A point is in shadow if it is **not** visible from the light source



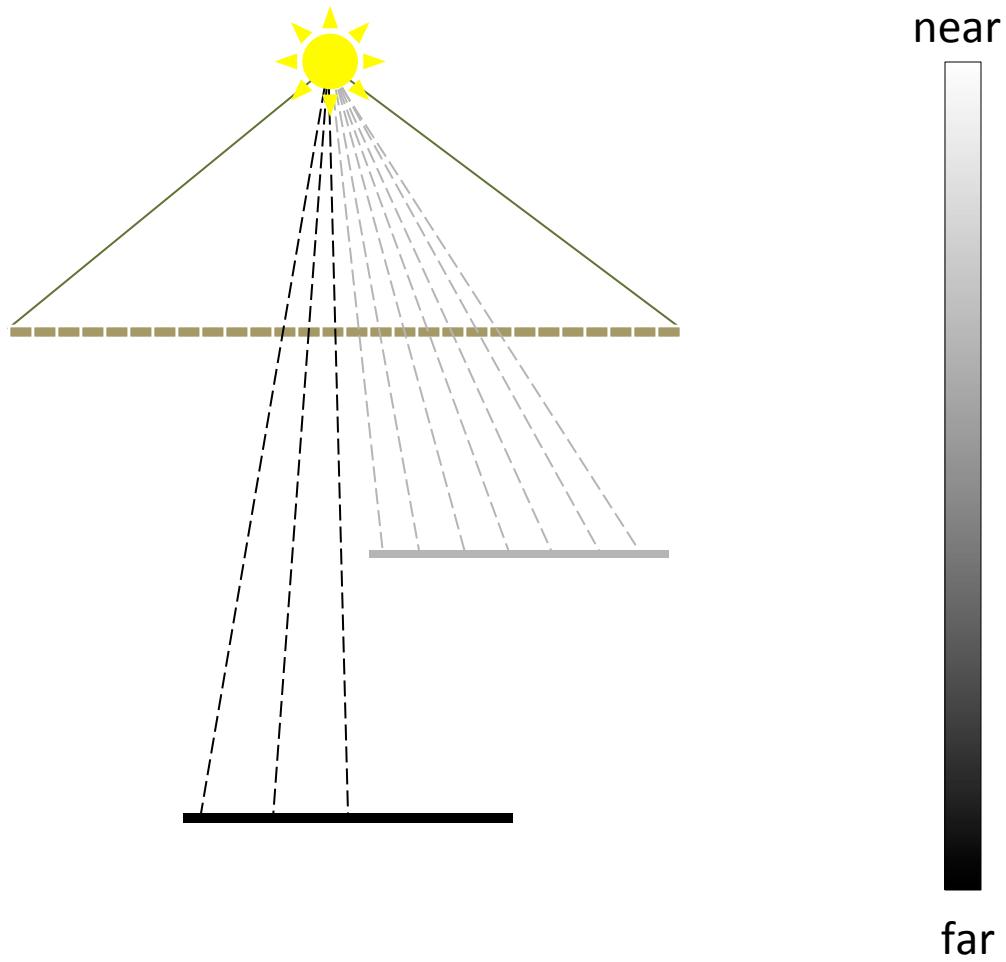
# Two passes

- Shadow map pass
  - Use the light source as a view point (light space)
  - Render scene
  - Store depth information in a shadow  $z$ -buffer = **shadow map**
- Shading pass
  - Render scene as usual form the cameras view point
  - Each pixel's position  $(x_v, y_v, z_v)$  is transformed to light space  $(x_s, y_s, z_s)$ ,
  - If
    - the  $z_s$  value is less or equal to the shadow map at  $x_s, y_s$  it is **lit**,
    - else it is **shadowed**

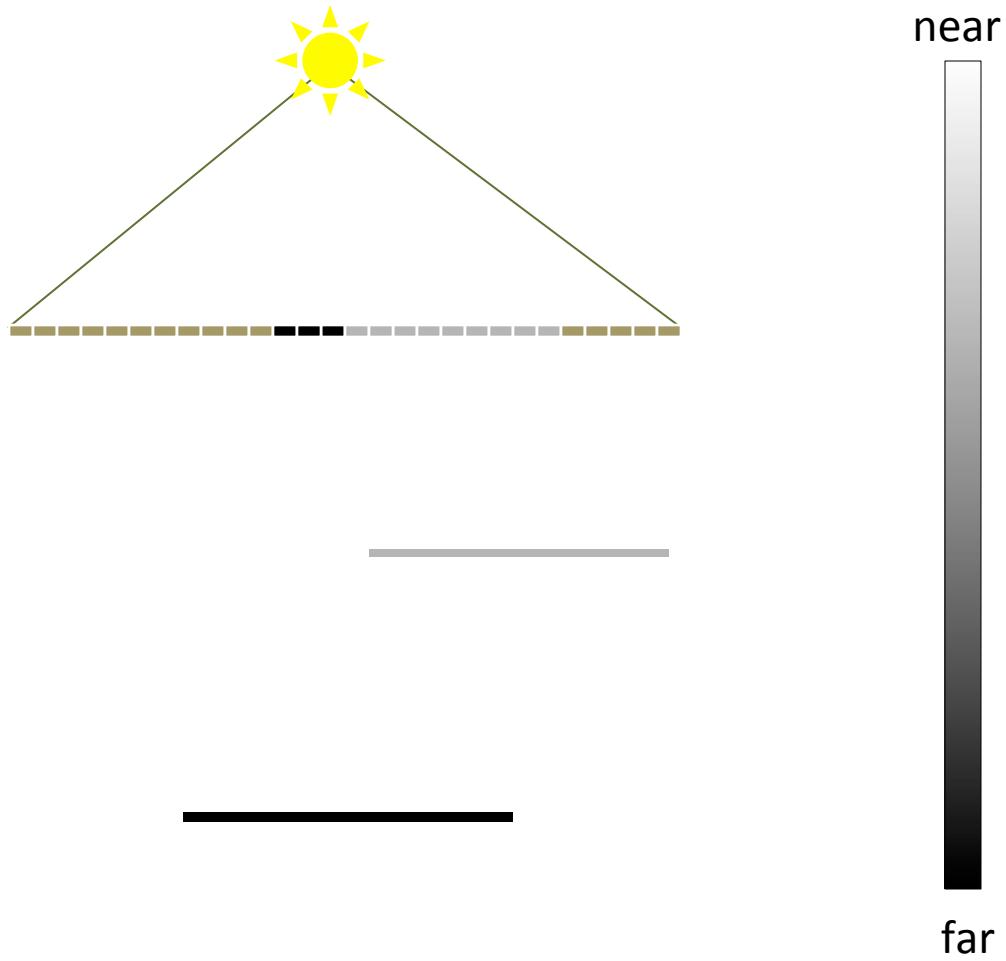
# Shadow Maps



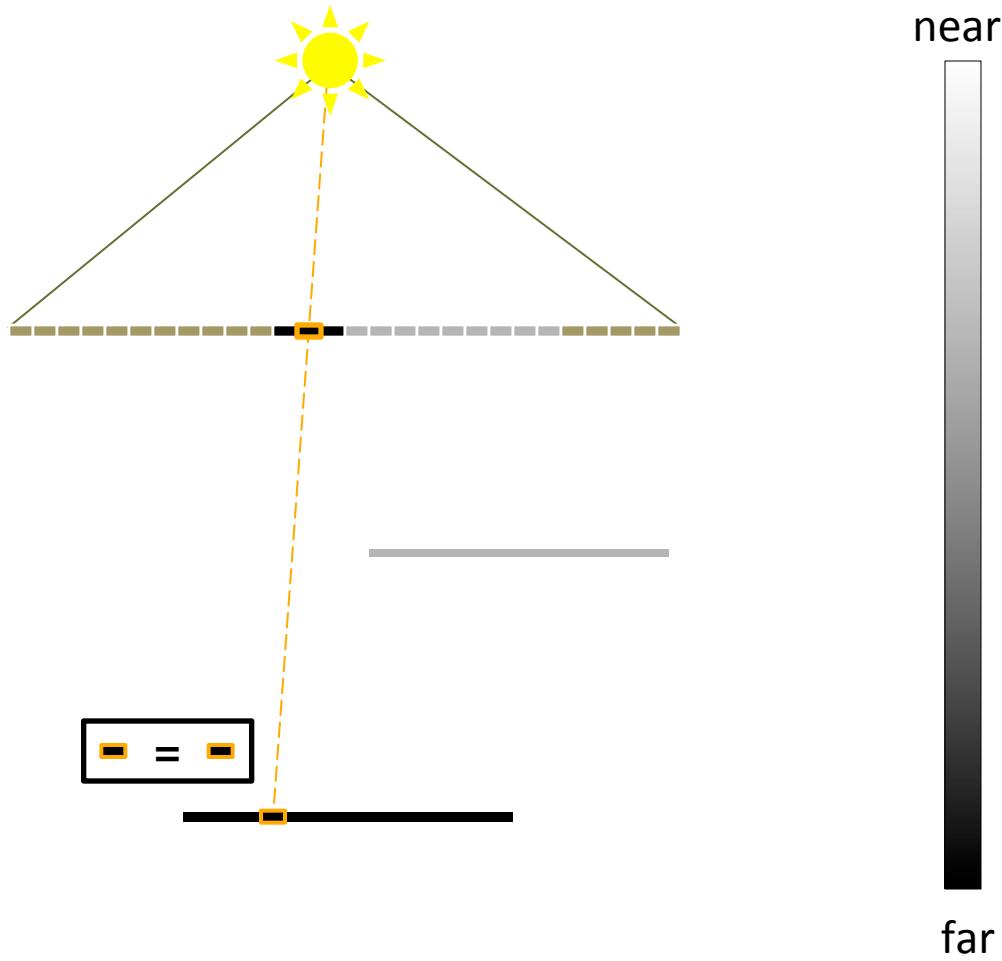
# Shadow Maps



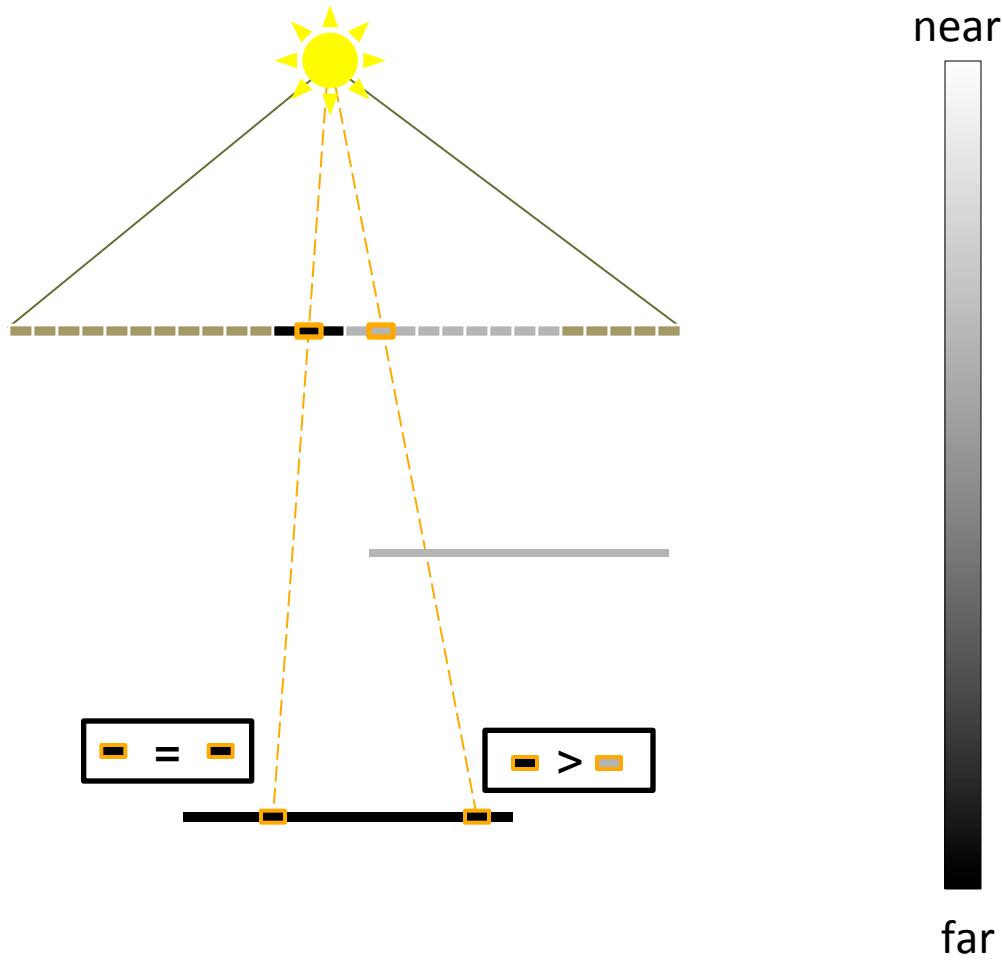
# Shadow Maps



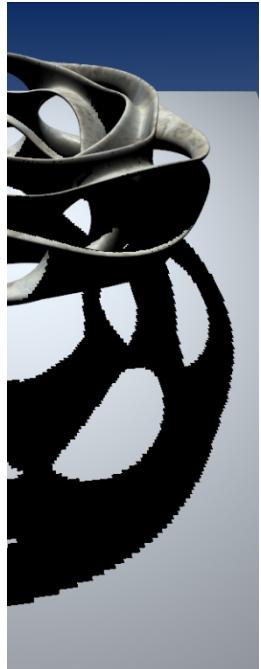
# Shadow Maps



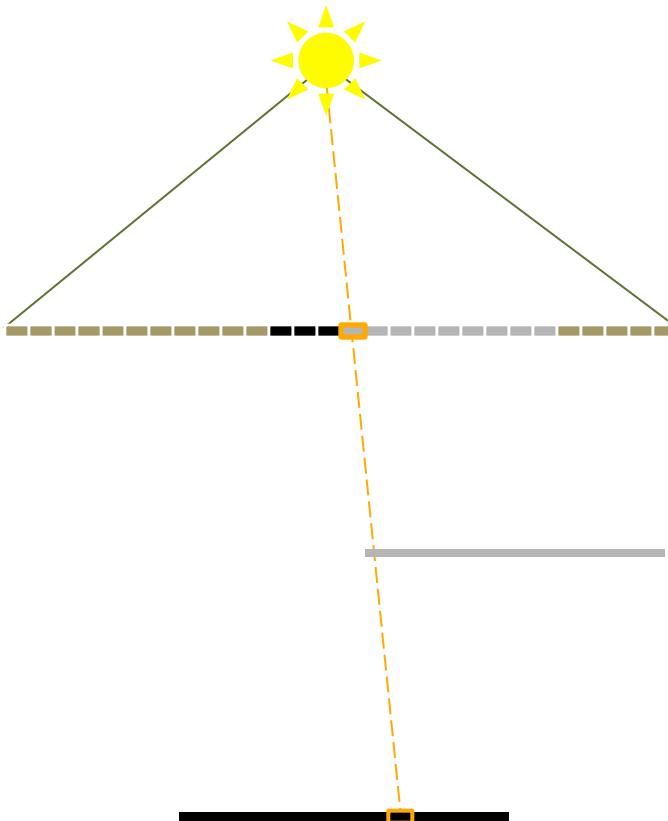
# Shadow Maps



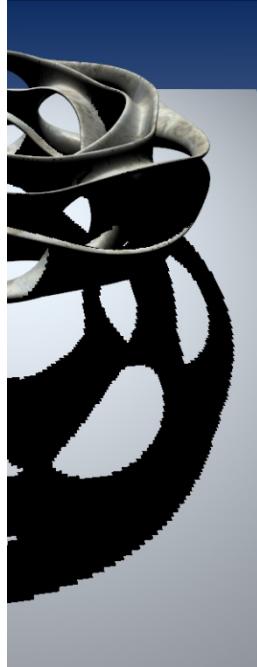
# Shadow Map Filtering



unfiltered



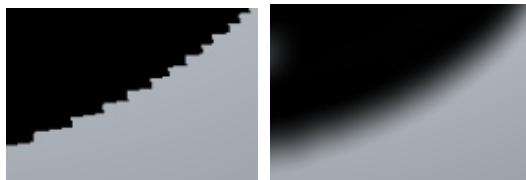
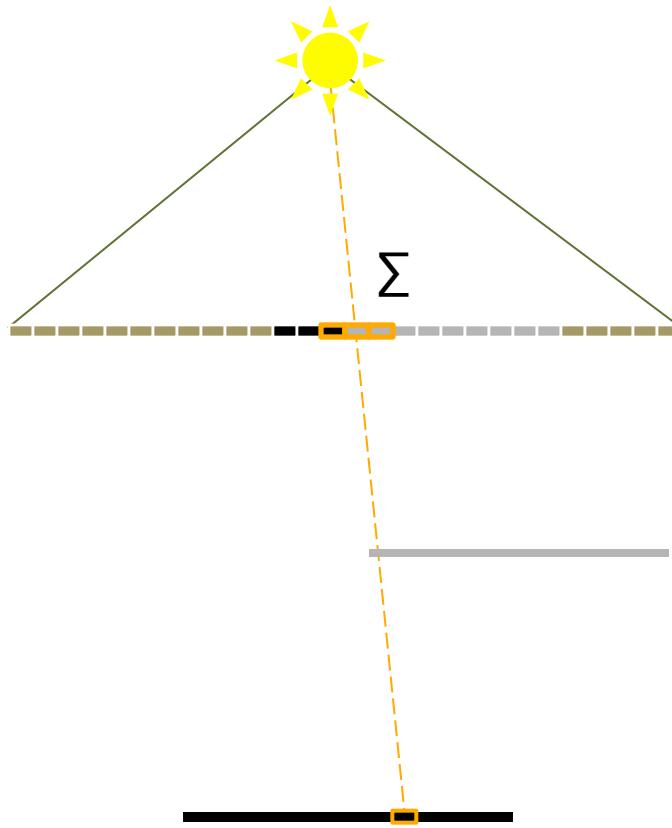
# Shadow Map Filtering



unfiltered



filtered



Percentage Closer Filtering  
[Reeves et al. '87]

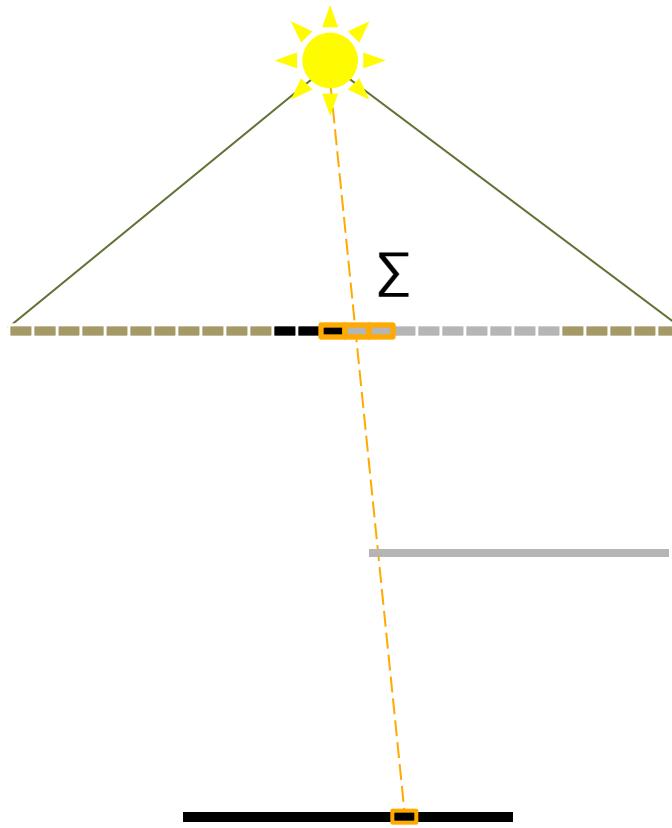
# Shadow Map Filtering



unfiltered



filtered



Percentage Closer Filtering  
[Reeves et al. '87]

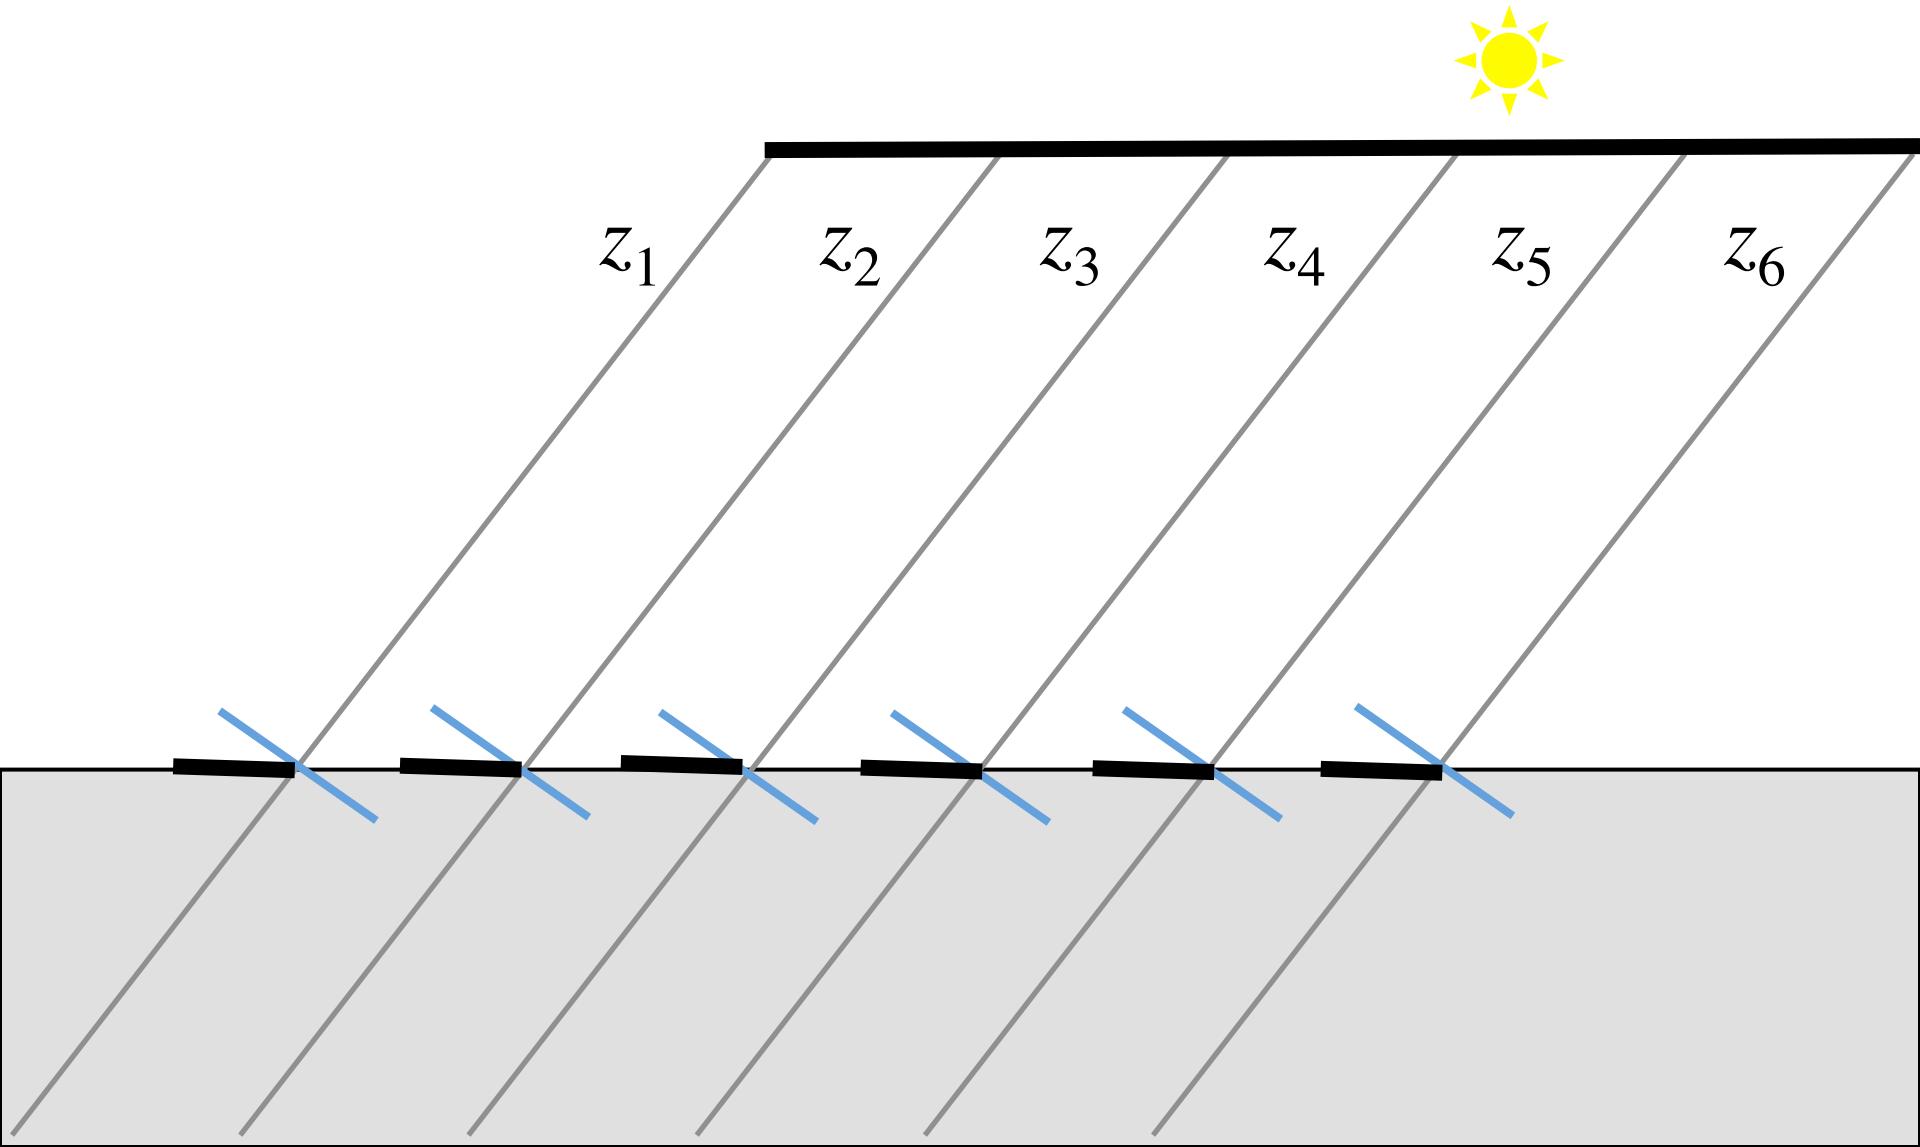
	1
	0
	0

compare first, then filter!

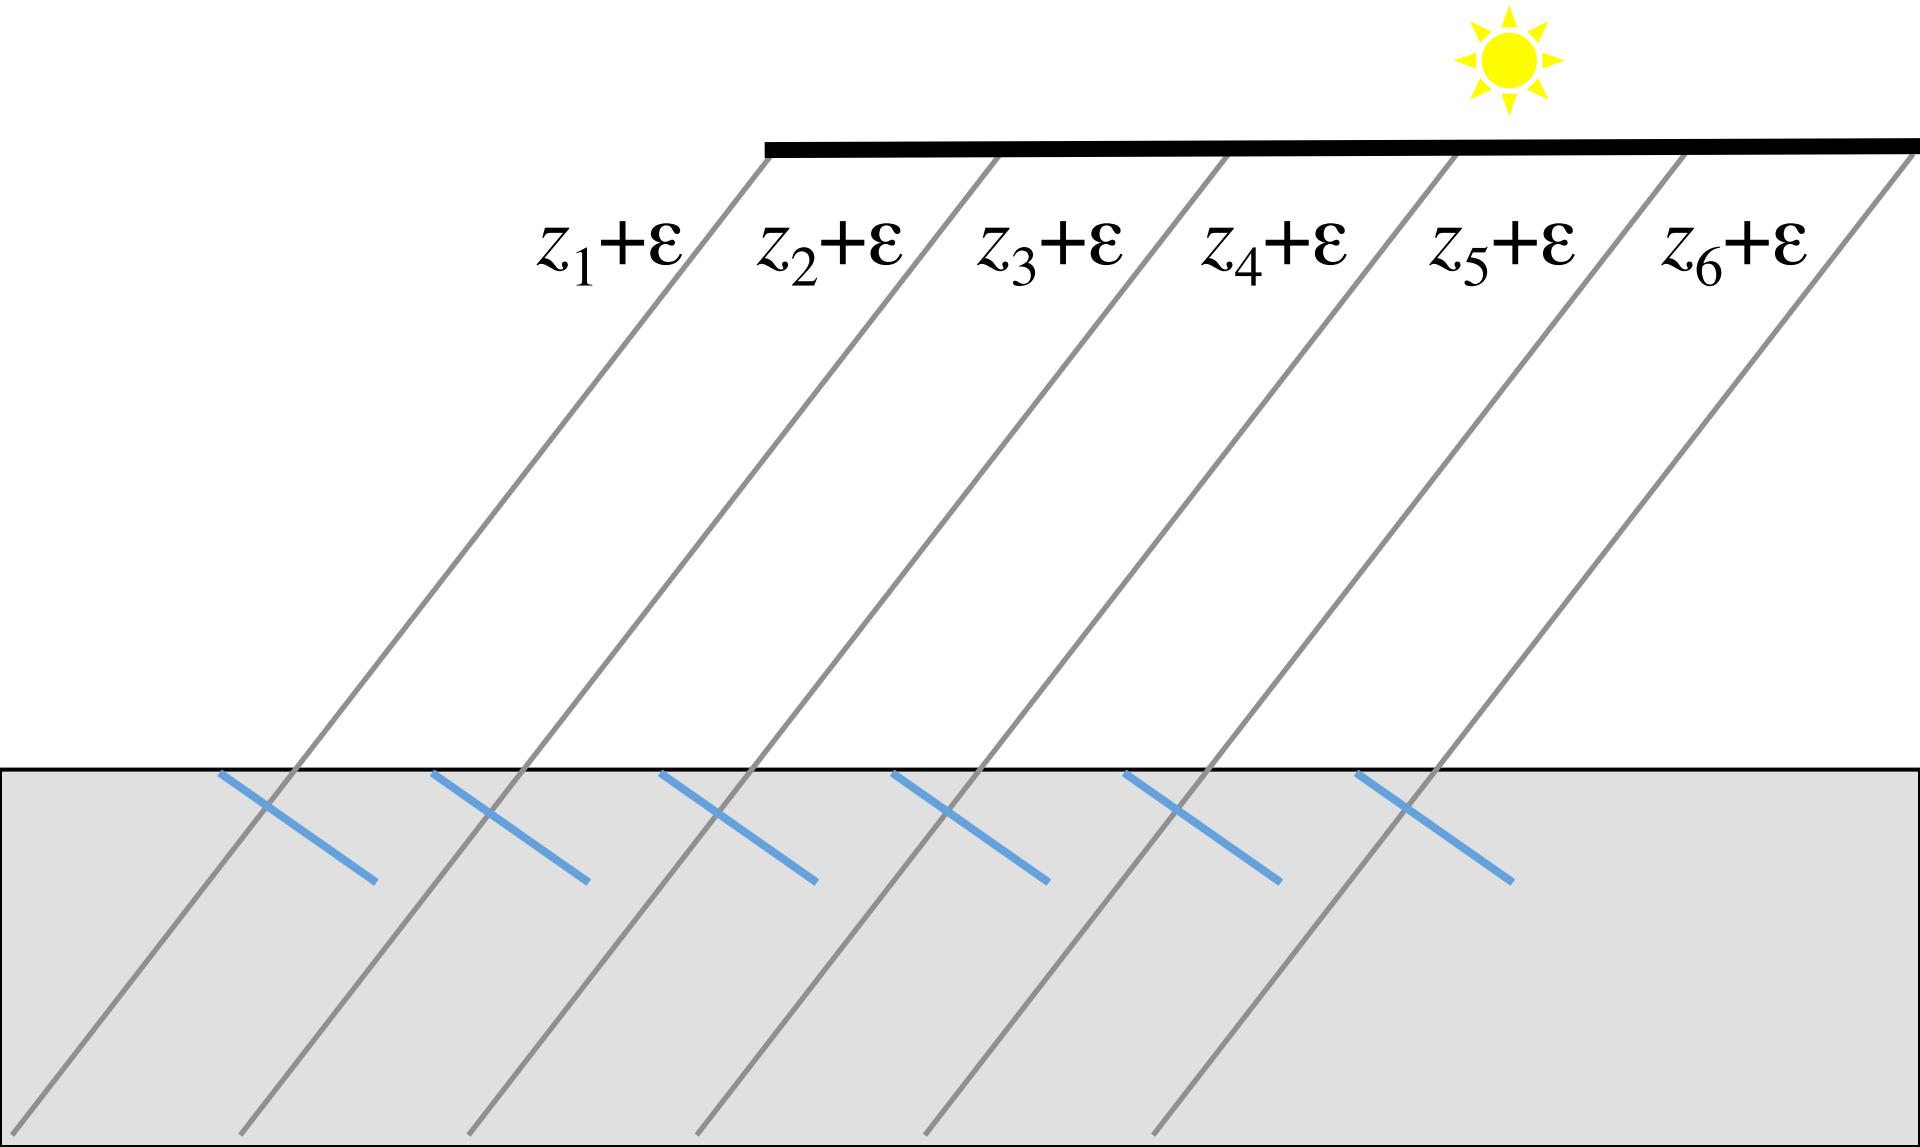


**Expensive**

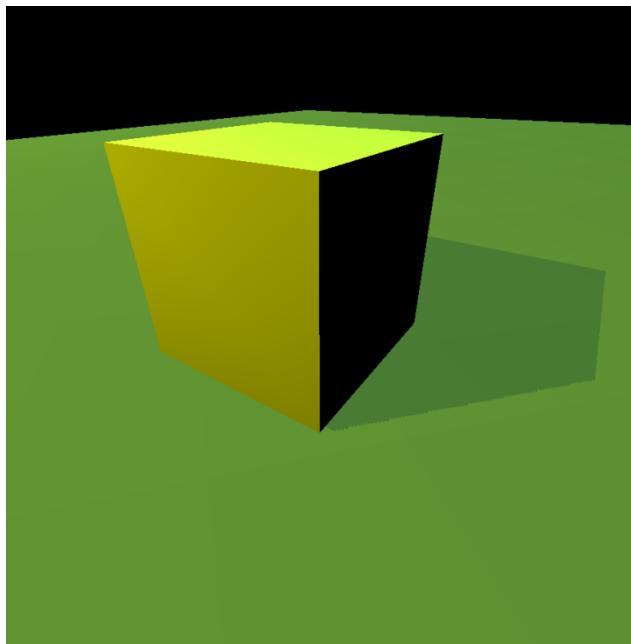
# Self shadowing: Disaster



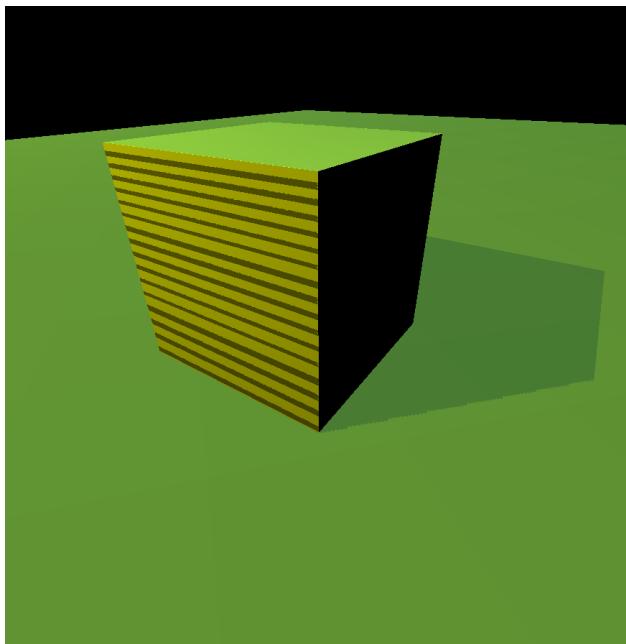
# Self shadowing fix



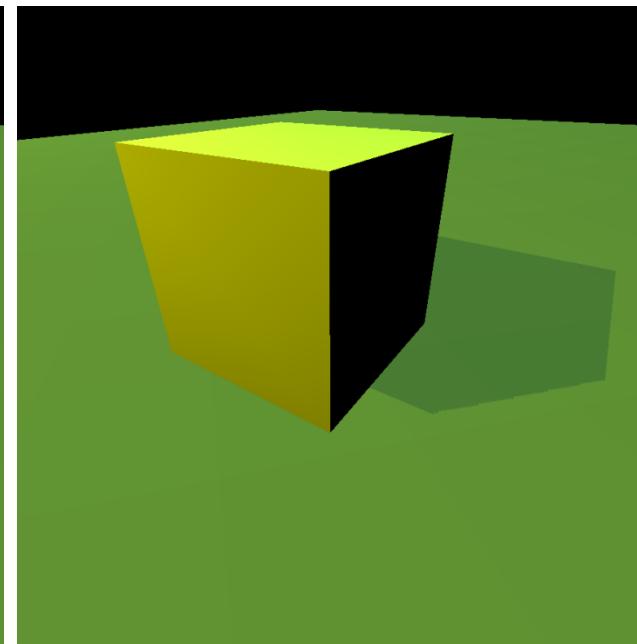
# Bias (Epsilon) for Shadow Maps



Correct result



Not enough bias



Too much bias

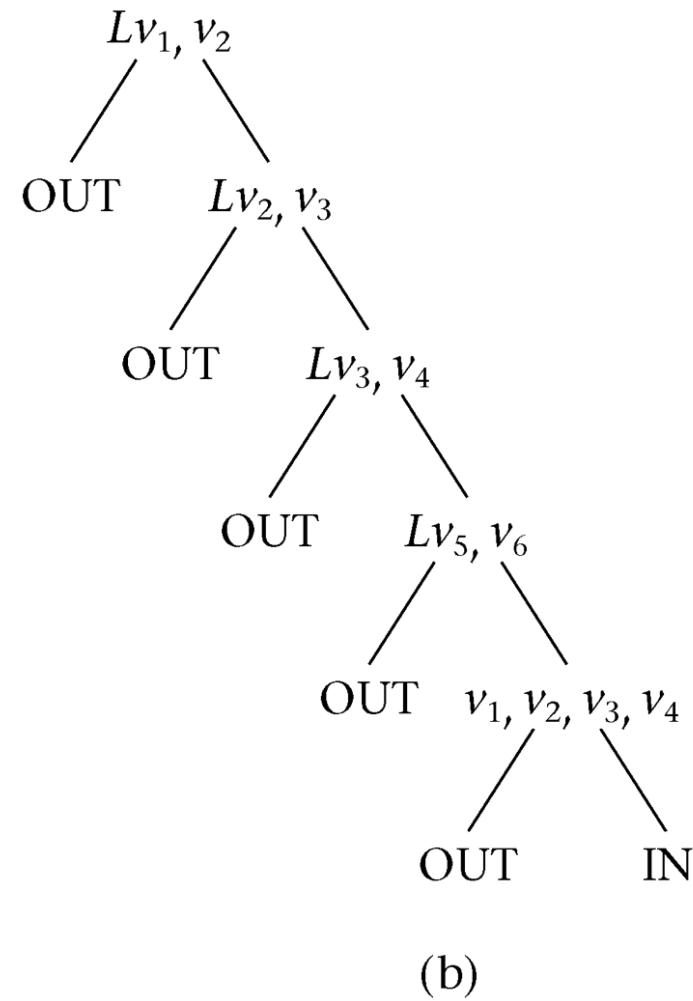
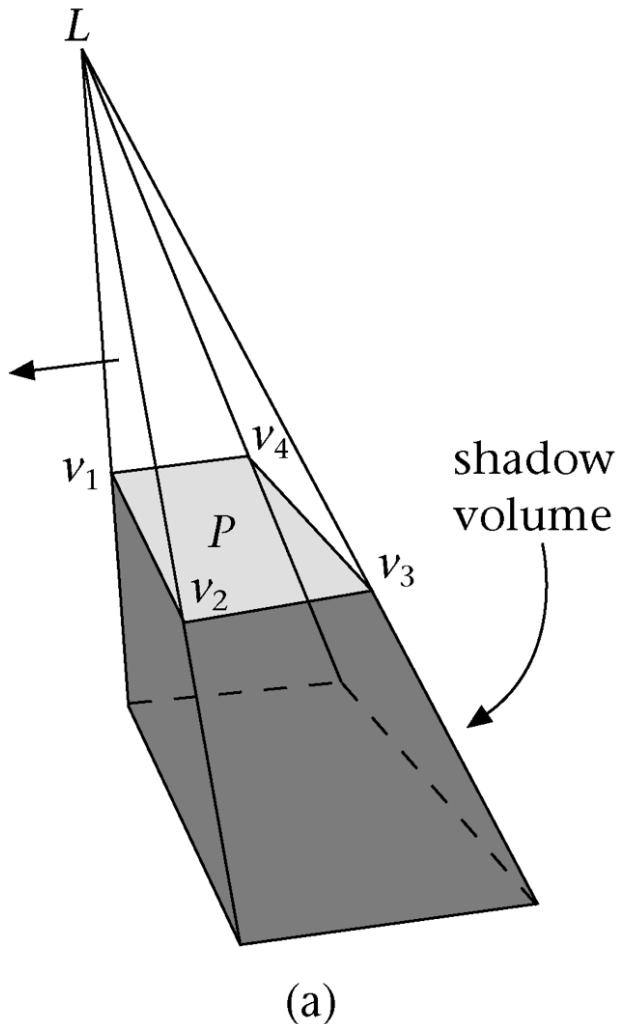
# Shadow Maps

- “Less than or equal” test is imprecise
  - Gives rise to “shadow acne”
- Often found in hardware now
  - Otherwise high-cost operation
- Imprecise since it is only accurate in the image space of the light
  - Imagine a shadow throw over complex objects or long distances
- Quality depends on resolution (jagged edges)
  - Percentage-closer filtering helps
- FOV of shadow map?

# Shadow Volume Method

- Shadow volume (SV) is the volume of space below a polygon that cannot see the source (a culled pyramid)
- During rendering of image, the line from a point visible through a pixel to the eye is intersected with all object SVs
- The number of intersections indicates if the point is in shadow or not

# Shadow Volumes



# Shadow Volumes

- Just like a polygon - you are inside a volume if you need to cross a surface to exit it
- General idea of shadow volumes is count the number of shadow planes you cross
  - +1 for front facing
  - -1 for back facing
- If total is  $>0$  you are in shadow
- Special case if the eye itself is in shadow

# Shadow Volumes

Two stages:

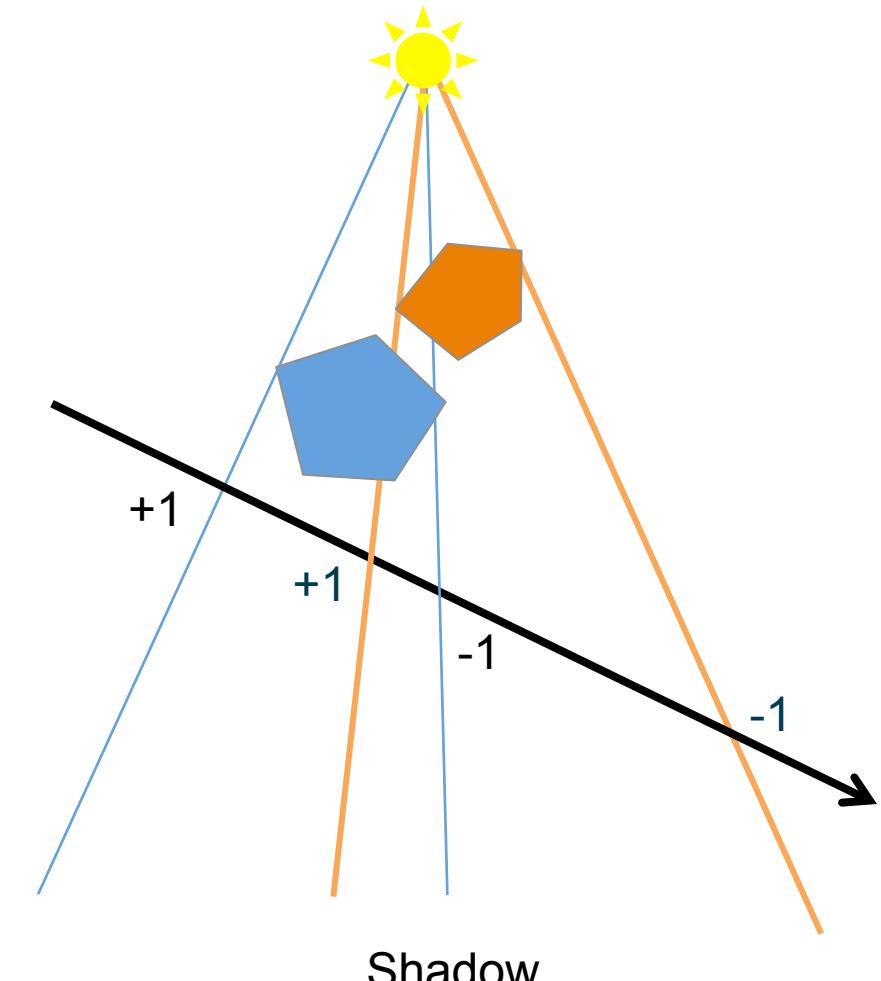
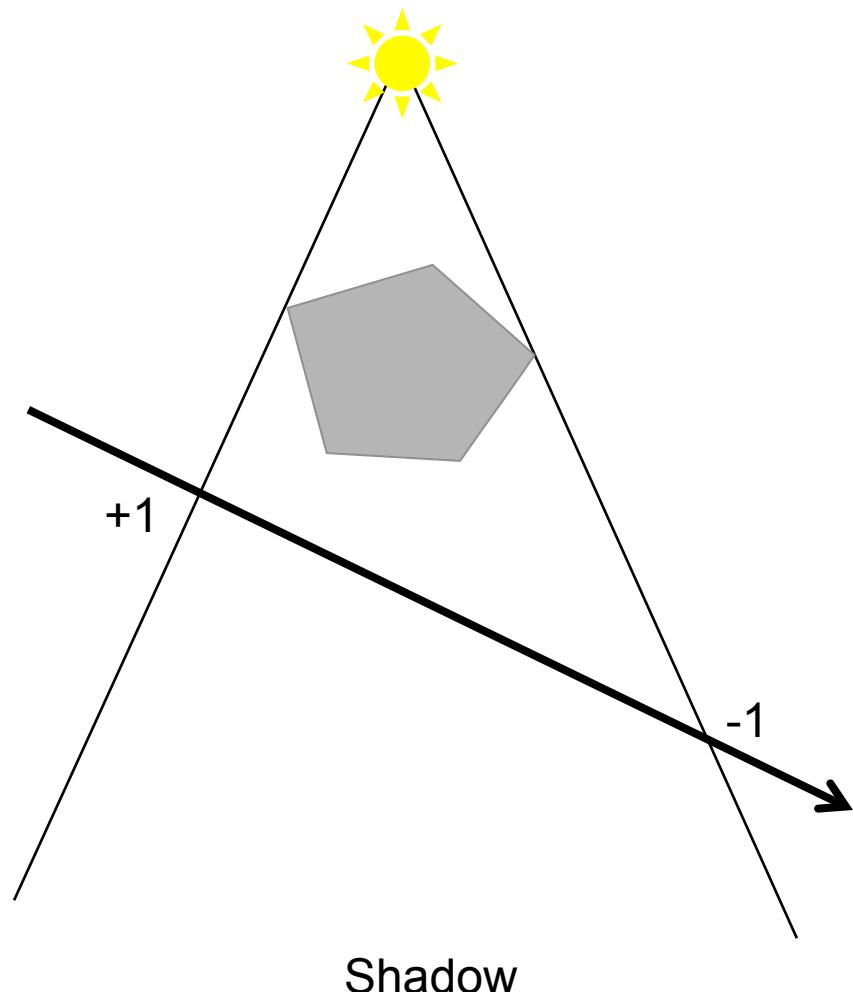
## 1) Volume **construction**

- Find all planes of the shadow volume and their plane equations

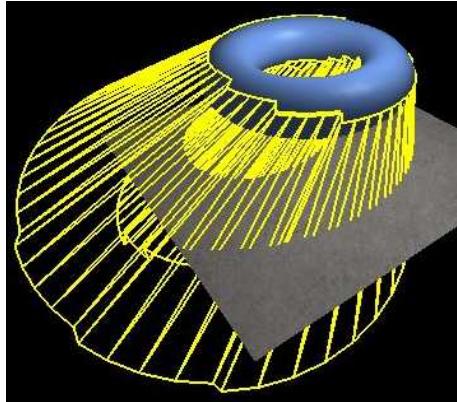
## 2) Volume **test**

- Determine shadow plane count per pixel
- Use a scan-line method OR stencil test

# Shadow Volume Example



# Shadow Volumes with OpenGL



- Shadow volumes are rendered at each frame
- The stencil buffer is used for counting how many SV are crossed
- Sometimes not all objects are used for casting shadows



# Shadow Volumes & Stencil Test

- A stencil buffer is screen sized buffer (1-8 bit) that stores a flag about a rendering operation
  - E.g.,  $\text{stencil}[x, y]$  is negated if  $\text{zbuffer}[x, y]$  is less than current z value (i.e. stencil is set if and only if z buffer-test passes)
- Many uses in graphics

# Shadow Volumes & Stencil Test

- Render the scene into the RGB and  $z$ -buffer
- Turn  $z$ -buffer **writing** off
- Render all shadow polygons with the stencil buffer
  - Increment stencil count for front-facing
  - Decrement for back facing
- Re-render scene with lighting **off** and only render pixels where stencil is non-zero

# Summary for Sharp Shadows

- Four shadow umbra techniques
- Image space
  - Shadow maps
  - Shadow volumes
- Object space
  - Fake shadows

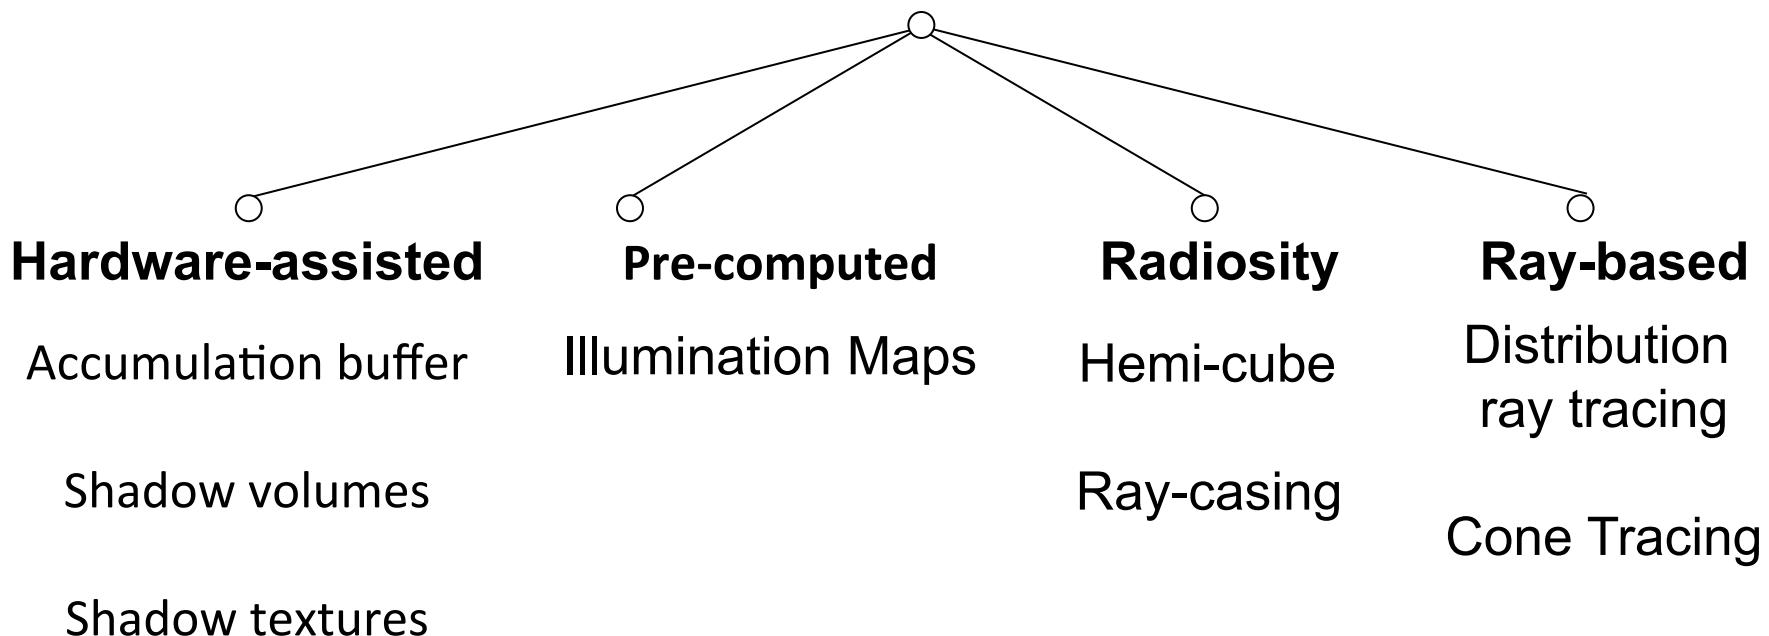
# Soft Shadows

# Soft Shadows

- Source has a finite extend
- Images look a lot more realistic



# Soft Shadows

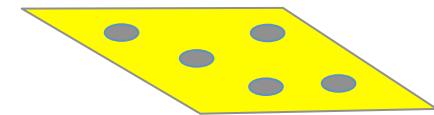


# Analytical v. Sampling

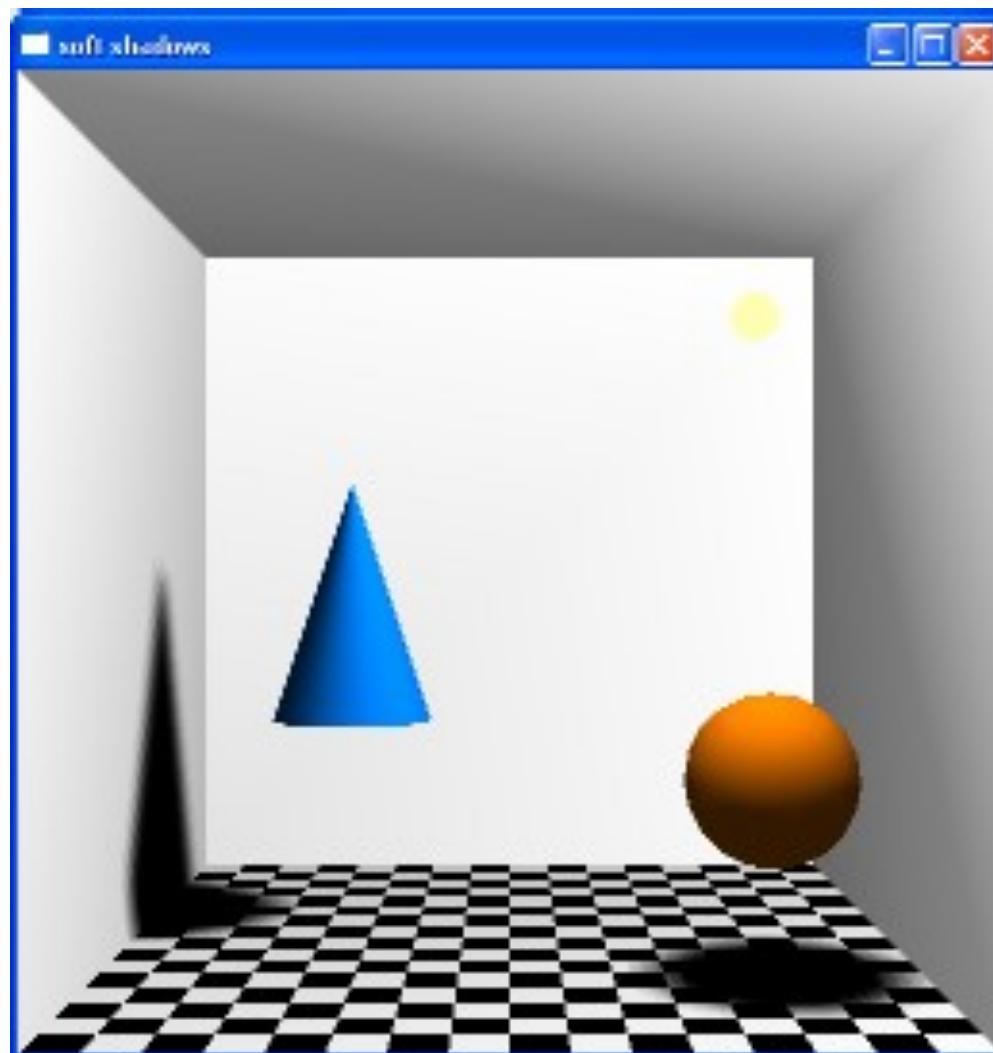
- Analytical
  - Find all boundaries within the penumbra.  
Done almost exclusively for polygonal light sources
- Sampling
  - Approximate solution that treat the light source as a set of points. Any shape source is possible.

# Soft Shadows using Point Sampling

- Place many point lights on an area light
  - Random positions work just fine
- Render hard shadows from each point light
  - E.g., using shadow volumes or shadow maps
- Sum up all contributions
  - Can be done on the GPU (in the frame-buffer)
- Similar to what ray-tracing does to get soft shadows



# Example



# Illumination Maps (Shadow Textures)

- Shadows are pre-computed and stored as textures on the receiving polygons
- Displayed using graphics hardware in real-time
- Often use: Radiosity
- Sometimes called “baked” lighting, very common in game engines
- Disadvantage: lighting cannot change

# Reacp

- In order to regain shadows in real-time engines, we have to do a lot of work
- A very large number of shadow algorithms exist
- Many of them are unsuitable for walkthroughs of very complex scenes:
  - with pre-computation methods scene cannot be modified
  - or are too slow (ray-tracing, soft shadows)
- Hard shadows
  - on-the-fly methods (SM and SV) are fast enough