

# Monte-Carlo Path tracing

# Today

- Solving integrals (approximately)
  - What did ray-tracing do?
  - Analytic vs. Numeric
  - Monte Carlo method
- Solving the RE using Monte Carlo
- Variance reduction
  - Good Sampling patterns
  - Importance sampling

# Why is this hard to solve?

- It involves an integral with no analytic solution
- It is an integral equation, so the RHS contains the LHS in an integrand

(Spherical) integration

$$L(\mathbf{x}, \omega) = L_e(\mathbf{x}, \omega) + \int f_r(\mathbf{x}, \omega_i, \omega) L(\mathbf{y}, -\omega_i) \cos \theta_i d\omega_i$$

Same thing!

# How ray-tracing solves the RE

- Many got suspicious about ray-tracing
- Example: Does metal have finite gloss?
  - Yes, cause otherwise I cant see highlight!
  - No, reflections are not blurry in steel balls!
- Contradiction!



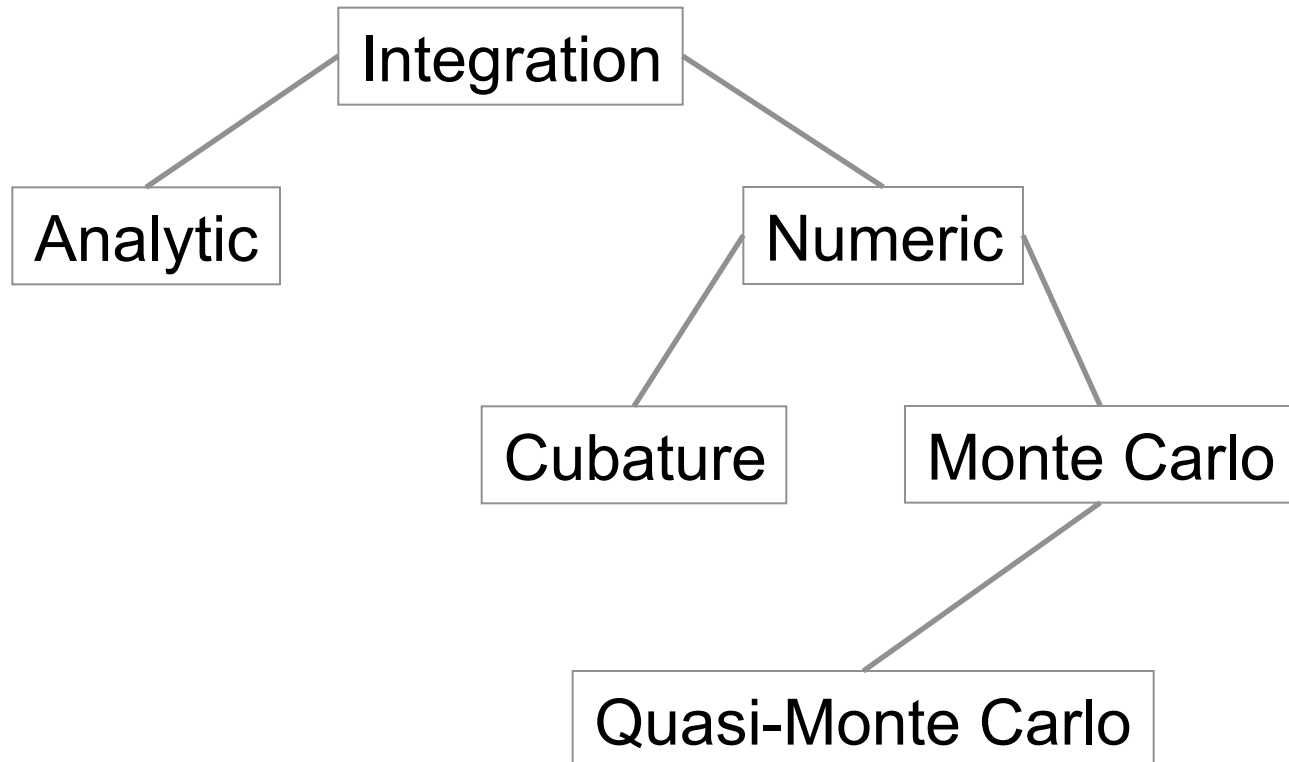
# How ray-tracing solves the RE

- The integral is split into a sum of two:
  - A Dirac-paths, solved by binary recusing
  - A path connecting to a point light, can evaluate without recursion

# Solving integrals

- Analytic (accurate)
  - Given the function  $f$ , try to find  $F$  by symbolic manipulation
- Numeric (approximate)
  - Cubature
  - The Monte Carlo method
  - Both are approximate

# Classification



# Analytic

- Examples

- $f(x) = x$ ,  $F(x) = \frac{1}{2} x^2$

- $f(x) = \sin(x)$ ,  $F(x) = -\cos(x)$

- Difficult for a function such as we have


- Impossible, as: The input is not even analytic (what is  $f$  for a bunny in the sun?)

- Difficult, as: Recursion

- Also difficult: Spherical domain



# Some analytic things work




$$f(\boldsymbol{\omega}) = f(\theta, \varphi) = \begin{cases} c & \text{if } 2 < \theta < 3 \text{ and } -1 < \varphi < 1 \\ 0 & \text{otherwise} \end{cases}$$

# Forget about spheres for now

Integrand

$$f(\omega)$$

Dark = low value  
Bright = high value




Domain

$$\Omega$$


(Here 2D rectangle)

$$\text{Solution } F = \int f(\omega) d\omega = 0.85$$

# Numerical integration




# Numerical integration




You should recall from high school analysis course.

# Numerical integration

$1 + 1 + 1 + 1 +$   
 $0.2 + 0.2 + 1 + 1 +$   
 $1 + 0.2 + 0.7 + 1 +$   
 $1 + 1 + 1 + 1$   
 $\rightarrow$   
 $13.3 / 16 = 0.83125$   
 $\sim$   
 $0.85$




# Numerical integration



# Numerical integration

...  
→  
 $54.2 / 64 = 0.846$   
~  
0.85



# Cubature: A bit more formal

$$F(\mathbf{w}) \text{ on } \Omega = \int f(\mathbf{w}) = \Sigma f(\mathbf{w}_i) / N$$

To find the integral  $F$  of a function  $f$ , decompose  $\Omega$  it into  $N$  as-small-as-possible cubes, evaluate  $f$  on each (**sample**) and average.



# Is this it?


- Very simple method!
  - To get  $L(\mathbf{x}, \omega_o)$
  - Subdivide hemisphere  $\Omega$  above  $\mathbf{x}$  into  $N$  strata  $\omega_i$
  - Evaluate integrand, i.e., send a ray
$$f = L(\mathbf{y}, -\omega_i) f_r(\mathbf{x}, \omega_i, \omega_o) \cos(\theta)$$
    - Triple product of light, svBRDF and geometric term
  - Average
  - Done!

# Problem: Curse of dimensionality

- $L$  appears on both sides
  - For  $L(\mathbf{y}, -\boldsymbol{\omega}_i)$  need to solve another integral
  - OK, let's go to  $\mathbf{y}$  and also compute  $L(\mathbf{y}, -\boldsymbol{\omega}_i)$ 
    - For  $L(\mathbf{z}, -\boldsymbol{\omega}_i)$  need to solve another integral
    - OK, let's go to  $\mathbf{z}$  and also compute  $L(\mathbf{z}, -\boldsymbol{\omega}_i)$ 
      - For  $L(\mathbf{z}_2, -\boldsymbol{\omega}_i)$  need to solve another integral
      - OK, let's go to  $\mathbf{z}_2$  and also compute  $L(\mathbf{z}_2, -\boldsymbol{\omega}_i)$ 
        - » For  $L(\mathbf{z}_3, -\boldsymbol{\omega}_i)$  need to solve another integral
        - » OK, let's go to  $\mathbf{z}_3$  and also compute  $L(\mathbf{z}_3, -\boldsymbol{\omega}_i)$ 
          - For  $L(\mathbf{z}_2, -\boldsymbol{\omega}_i)$  need to solve another integral
          - OK, let's go to  $\mathbf{z}_2$  and also compute  $L(\mathbf{z}_2, -\boldsymbol{\omega}_i)$

# Recursion

- Recursion in CS is not evil
- Worked well in ray-tracing:




For depth  $d$  and reflection/refraction the number of rays is

$$2^d$$

So three-bounce is  
 $2^3 = 8$  rays / pixel

# Recursion

- Recursion in CS is not evil
- Impractical for the real RE
- Typical  $N$  is maybe 100 to 1000




For depth  $d$  and  $N$  strata  
the number of rays is

$$N^d$$

So three-bounce is  
 $1000^3 = 1\text{B rays / pixel}$


# Another way to imagine it



Every time we want any accurate value (orange dot) at any bounce we need to resolve exponentially many others below it to proceed


# Alternative: Random samples

$$\begin{aligned} &11 * 1.0 + \\ &3 * 0.2 + \\ &2 * 0.7 \\ \rightarrow & \\ &13.0 / 16 = 0.8125 \\ &\sim \\ &0.85 \end{aligned}$$




# Alternative: Random samples

$11 * 1.0 +$   
 $3 * 0.2 +$   
 $3 * 0.7$   
 $\rightarrow$   
 $13.7 / 16 = 0.856$   
 $\sim$   
 $0.85$



# What do we get from random?




For depth  $d$  and  $N$  samples  
the number of rays is

$$N \times d$$

So three-bounce is  
 $3 \times 1000 = 3\text{K}$  rays / pixel



# Another way to imagine it



As we just need an approximate value, we can proceed with any point without looking at all others

# Monte Carlo: A bit more formal

$$F(\omega) \text{ on } \Omega = \int f(\omega) = \Sigma f(\omega_i) / N$$


To find the integral  $F$  of a function  $f$ , place as many samples  $N$  onto  $\Omega$ , evaluate  $f$  on each and average.

# This is it

- Still very simple method!
  - To get  $L(\mathbf{x}, \boldsymbol{\omega}_o)$
  - $N$  times
    - random directions  $\boldsymbol{\omega}_{i,0}$  above  $\mathbf{x}$ ,  $\boldsymbol{\omega}_{i,1}$  above  $\mathbf{y}$ , etc.
    - Evaluate integrand, i.e. send a ray
$$f = L(\mathbf{y}, -\boldsymbol{\omega}_i) f_r(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o) \cos(\theta)$$
      - Triple product of light, svBRDF and geometric term
  - Done!

# How to pick random directions?

- How to pick a random ray in 3D?
- `normalize(vec3(frand(), frand(), frand()))?`
- Clumps on axis and diagonals
- Bias in result!



# When to stop?

- Multiple options
- Popular:
  - After a fixed depth
  - When contribution falls below a threshold

# Progressiveness



1 Sample


# Progressiveness



2 Samples




# Progressiveness



10 Samples



# Progressiveness



1000 Samples

# Progressiveness

- After  $N$  samples we get an image
- After  $2N$  samples, we get an even better one
- This is called **progressive**
- Very useful for previews

# Random has another reason

- **Aliasing** is the second reason for random
- Consider this integrand, not even recursive:

Cubature with 1 sample = 0.0



Cubature with 2 samples = 0.0



Cubature with 5 samples = 0.0



Ground truth = 0.5

# Random has another reasons

- **Aliasing** is the second reason for random
- Consider this integrand, not even recursive:

Monte Carlo with 1 sample = 0.0



Monte Carlo with 2 samples = 0.5



Monte Carlo with 5 samples = 0.4




Ground truth = 0.5


# Estimator/Variance/Bias

- We get a new value for every random seed
- We map these into an **estimate**
- This is the value of the integral
- If there is a deviation, we call it **bias**
- Around this exists a distribution of values
- This distribution has a **variance**

# Variance of an estimator



# Bias of an estimator



# Desiderata variance


- Also it has no bias, i.e., the expected value is really the solution of the integrand
- OpenGL and CW1 ray-tracing: All biased
- A good estimator has a low **variance**
- Whenever we render, we will get a value close to the true value
- To this end, we do **variance reduction**



# Variance reduction


- Next-event estimation
- Good sample patterns
  - Jittered
  - Quasi-Monte Carlo
  - Blue noise
- Importance Sampling

# Path tracing - High hopes




We hope for this ...

# Path tracing – High hopes




We hope for this ...

# Path tracing - Reality




But what we get is

# Path Tracing - reality




But what we get is

# Next-event estimation



But what we get is

# Problem: Double-accounting



There are now **two ways** to hit the light.  
Simple solution: Simply only take emission from NEE.


# Next event estimation

- Two simple changes
  - Add a random ray in the direction to the light
  - Remove adding in emission  $L_e$  on all other paths
- Best for **small** light sources
- Result:
  - Will never miss direct light at any point
  - Still have all benefits of MC
- Glorified Whitted-style CW 1 ray-tracing




# Endpoint choice

- Simple for flat area lights




# Endpoint choice

- Hard for e.g., spheres



# Uniform random can go bad

16  
→  
 $16 / 16 = 1$   
~  
0.85  
(not so cool)




# The ideal sample pattern

- Two contradicting goals:
  1. Maybe not be regular
  2. But always cover the domain uniformly, not only in the limit


# Jittering

- Back to the future:
  - Do cubature first
  - Then jitter every sample inside its cell
- Suffers from the curse of dimensionality
- Prevents aliasing
- Applicable if dimensionality is low
  - Example: Area light sampling


# Regular (recap)



# Jittered



# Multi-Jittered







# Halton

- A way to place samples
  - Somewhat uniform
  - Without structure
  - In high dimensions
- A typical work-horse solution for rendering
- Defined on the unit hypercube


# Radical inverse base $n = 2$





# Radical inverse base $n = 3$



Halton in 2, 3 i.e.  $\pi(0)$ ,  $\pi(1)$



# Random vs. Halton




# Quasi-Monte Carlo sampling

- To get an  $n$ -dimensional Halton pattern
- Build the radical inverse in the co-prime basis  $\pi(0, \dots, n-1)$
- Build tuples in the order they occur in each sequence
- Improvement: **Hammersley** (regular 1<sup>st</sup> dim.)
- De-correlation: **Cranely pattern** rotation

# Poisson disk / Blue noise

- Patterns with a maximal minimal distance between all points is called **Poisson disk**
- Another way is to see the spectrum of the distribution of distances: It is blue, i.e. no small minimal distances


# Poisson disk



*Note:* The smallest circle of all circles around each point is quite large




# Blue noise



Receptor distribution on the macaque retina prevents aliasing

# Importance sampling



Put more samples where the integrand is high, as here the errors have the largest effect.

# Importance sampling

- Placing the samples non-uniformly will introduce bias
- Fortunately, random uniform is just a special case of a more general estimator formulation we will see next
  - Before we took  $\omega_i$  uniform, so  $p(\omega_i) = 1 / |\Omega|$
  - Any other  $p$  will work as well
  - Ideally  $p \sim f$

# MC with importance sampling

$$F(\omega) \text{ on } \Omega = \int f(\omega) = 1/N \sum f(\omega_i) / p(\omega_i)$$

To find the integral  $F$  of a function  $f$ , place as many samples  $N$  onto  $\Omega$  according to a distribution  $p$ , evaluate  $f$  and  $p$  on each and divide.

# What can be $p$ ?

- Recall the integrand is
$$f = L(\mathbf{y}, -\boldsymbol{\omega}_i) f_r(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o) \cos(\theta)$$
- We could sample for
  - Light: Hard, integral equation itself.
  - BRDF: Not so hard, done analytically
  - Geometric term: Even easier analytically
  - Products of all of the above: Even harder than any alone, but doable

# Example: IS for direct light



Same amount of rays ©U Virginia

# Recap

- Rendering is solving an integral equation
- Analytic and some numeric methods no-go
- Monte Carlo is the method of choice
- Suffers from noise (variance)
- Need to use variance reduction methods
  - Sample patterns
  - Importance sampling